




INVESTMENT RECOMMENDATION · SEPTEMBER 21, 2026



Exceptional Business, Fully Priced — The Bear Case Bets Against Management's Own Guidance

Base Case

LEXO INVESTMENT RESEARCH

CONFIDENTIAL

NVIDIA Corporation — One-Page Summary

KEY FACTS

*All figures in this memo are as of 2026-09-08. Every rating, price target, stop-loss and return below is measured relative to the current price of **\$230.36** — the 2026-09-04 close, re-confirmed on 2026-09-08 (2026-09-07 was Labor Day, so there is no closing print between the two dates).*

FACT	VALUE	AS OF	SOURCE
Current Price	\$230.36	2026-09-04 close, re-confirmed 2026-09-08	Trading · Implied vs Current
Market Cap	\$5,551,676M	2026-09-08	$\$230.36 \times 24,100M$
Shares Outstanding	24,100M diluted	FY2026 DCF basis	Trading · DCF Mirror
52-Week Range	\$164.07 – \$236.54	2026-09-08	Trading · Price OHLC
Position in Range	2.6% below the high, 40.4% above the low	2026-09-08	Trading · Price OHLC
Base-Case Intrinsic (DCF)	\$139.24	2026-09-08	Trading · DCF Mirror / Implied vs Current
Probability-Weighted Intrinsic	\$136.78	2026-09-08	Trading · Implied vs Current
Implied vs Current	–40.6%	2026-09-08	Trading · Implied vs Current
Bear / Base / Bull Implied	\$78.83 / \$139.24 / \$189.80	2026-09-08	Trading · Implied vs Current
Enterprise Value (market-implied)	\$5,510,019M	2026-09-08	Trading · Implied vs Current
Dividend Yield	not sourced	—	Not carried in NVDA Reference or any research pod; NVIDIA pays a nominal dividend, but no yield is asserted or computed here
Next Earnings Date	2026-11-17 (Q3 FY2027, company-confirmed)	2026-09-08	Trading · Catalyst Timeline
Rating / Conviction / Position	Hold · Low · no position initiated	2026-09-08	Trading · Cover / Execution Plan
Stop-Loss (existing holders only)	\$195.00 on two consecutive weekly closes (–15.3% from \$230.36)	2026-09-08	Trading · Execution Plan

Market cap computed as current price × diluted shares outstanding ($\$230.36 \times 24,100M = \$5,551,676M$); market-implied enterprise value is that market cap less \$41,657M of net cash. The 52-week extremes are the quoted levels and differ by cents from this workbook's own monthly bars (\$236.26 high, \$164.08 low) — both are reported rather than reconciled. Next Earnings Date is the company-confirmed 2026-11-17 print from Catalyst_Timeline; Reference_Card carries the older, unconfirmed "estimated 2026-11" placeholder. Prices are a build-time snapshot as of 2026-09-08 and do not auto-refresh — regenerate to re-anchor.

At the **\$230.36** close on 2026-09-04 (re-confirmed 2026-09-08), NVIDIA is an N-of-1 business priced **40.6% above** this workbook's probability-weighted intrinsic value of **\$136.78**. The gap is a claim that management's own FY2028 guidance is wrong, and it is testable on **17 November 2026**. Until then: **Hold, Low conviction, no position initiated**.

Lexo Verdict (3–5yr intrinsic): Diamond in Air | Ocean | Luminous — **N-of-1, 83/100**, multibagger gate BORDERLINE; trajectory Bull 25% · Base 50% · Bear 25% (15% competitive, 10% geopolitical).

Trading Stance (6–12mo tactical): Hold · Low conviction · no position initiated · PWR **–13.1%** modelled, **–9.7%** with the \$195.00 stop applied — a 12-month price expectation, not the –40.6% intrinsic gap.

Reconciliation: The 3–5 year verdict is unchanged and positive, the 6–12 month stance is Hold, and they do not conflict — both horizons carry identical 25/50/25 weights, and the –40.6% gap localises to one row: FY2028 revenue of **\$474,927M (26.8% growth)** against consensus **\$681,500M** and management's guided **~70%**. The bet is against company guidance, not market exuberance: forward EV/EBITDA of 20.3x sits *below* NVIDIA's own five-year low of 40.0x, and a 24.7x forward P/E is a 12.8% premium to the 21.9x peer median. Resolution: **17 November 2026**.

Why we like it (or don't)

- **The quality question is closed.** 68.8% ROIC, a 74.7% TTM gross margin against a 58.4% peer median, a 65.2% EBIT margin against 21.5% — three of four Lexo axes at their ceiling.
- **Share loss reads as mix, not defeat.** ~11 points of accelerator share ceded since 2023 while gross margin *rose* 72.7% → 74.7%; networking compounds at 199% against 92% for the segment, ~\$39B a year incremental against ~\$26B foregone.
- **But every scenario fails at this price.** Bear/base/bull intrinsic of \$78.83 / \$139.24 / \$189.80 leaves even the bull **–17.6%** below \$230.36; reverse-solved, \$230.36 needs 6.95% perpetual growth or a 44.2% five-year revenue CAGR.

What could break the thesis

- **Our FY2028 is simply wrong:** the likeliest error in this workbook. FY2028 commentary holding near 70% with backlog detail on 17 November makes 20.3x genuinely low and turns this into a Buy.
- **Hyperscaler capex deceleration:** roughly half of Data Center revenue is four customers. Any two guiding capex flat or down in the late-January 2027 window, or FY2028 guidance below 45%, is the exit.

What we'd do

Hold at \$230.36; no new weight initiated. The Neutral-spring rule maps to a starter position and is deliberately overridden: this bag's own probability-weighted 12-month return is **–13.1%** against a **+2.7%** ceiling, so opening length would follow a rule against the analysis it serves. A reader already long stops at **\$195.00** on two consecutive weekly closes — **–15.3%** from \$230.36, beneath the \$199–\$201 shelf. Re-entry: FY2028 commentary near 70% with backlog detail on 17 November 2026. Horizon 6–12 months tactical.

All figures as of 2026-09-08 — see Key Facts above; evidence in the 5-page outline and memo body.

Lexo Spine

5-step analytical backbone — side view (time flows bottom → top) and top-down (looking along the time axis) on one page



5-STEP SPINE

- 1 Business Overview
- 2 Historical Analysis
- 3 Lexo Analysis
- 4 Trading Analysis
- 5 Investment Thesis

FACTOR CIRCLES

- Internal Factors
- Sector Factors
- Market Factors

FORCE TYPES


- Headwind
- Tailwind
- Wave +
- Wave -
- Existential Threat
- Mud

Side view: time flows bottom → top; each cone is its own column with a vertical spine — past hangs below the present plane, future rises above.

Top-down: looking along the time axis — bone is the central hub, factor rings around it.

Light Cone — NVDA

Concentric factor analysis: internal, sector & market forces



INTERNAL FACTORS

1	Reasoning-Era Inference Demand	Wave +
2	Annual Architecture Cadence	Wave +
3	Networking Attach Compounding	Wave +
4	Founder-Led Capital Discipline	Wave +
5	Inference Share Loss to Custom Silicon	Wave -
6	Earnings Quality Drift	Wave -
7	Opex Leverage Exhausted	Wave -
8	Key-Man Concentration	Wave -
9	Supply Ceiling on Volume	Mud
10	Self-Funded Demand Loop	Mud

SECTOR FACTORS

11	AI Infrastructure Buildout Wave	Wave +
12	Training-to-Inference Workload Shift	Wave -
13	Custom ASIC Substitution	Wave -
14	Memory and Foundry Cost Inflation	Wave -
15	Export Control Regime	Mud
16	Hyperscaler Capex Turns Debt-Funded	Wave -
17	Sovereign AI Demand Channel	Wave +
18	Sector Multiples Ignore Business Quality	Mud
19	Supply-Constrained Accelerator Pricing	Wave +
20	HBM and Advanced Packaging Scarcity	Headwind

MARKET FACTORS

21	Taiwan Strait Concentration	Existential Threat
22	Hyperscaler Capex Financing Shift	Wave -
23	Long-Duration Multiple Sensitivity	Headwind
24	US Export Control Regime	Mud
25	Two Diverging Chip Markets	Wave -
26	National-Interest Policy Protection	Tailwind
27	Rate Transmission Is Inverted	Tailwind
28	Conventional Macro Exposures Immaterial	Wave +

Thematic Forces — NVDA

Factors grouped by force type

28 Factors

10 positive · 13 negative · 5 other

WINDS Temporary forces (1–4 quarters)

4 factors

● HEADWIND (2)

20 HBM and Advanced Packaging Scarcity

HEADWIND

Sector

Memory and advanced-packaging capacity, not demand, is the binding volume constraint across every accelerator vendor — NVIDIA's CFO attributes the scarcity to the AI buildout itself, and the industry cannot add HBM or CoWoS capacity on the cadence demand is growing. Source: Pricing_Analysis.

IMPACT 0.80

23 Long-Duration Multiple Sensitivity

HEADWIND

Market

Market commentary now frames AI equities as trading like interest-rate-sensitive credit, and the anticipation of tightening is described as sufficient on its own — the expectation of hikes alone can compress valuations. This is a valuation channel rather than an earnings one, and it bites harder on a company whose value sits in distant cash flows. Source: Central_Bank_Watch.

IMPACT 0.75

● TAILWIND (2)

26 National-Interest Policy Protection

TAILWIND

Market

As of early 2026 no US regulatory action had successfully slowed NVIDIA, with its technology described as a national-interest priority in the race for AI supremacy. This rests on a single syndicated market-wire item and is an analyst's framing rather than a government designation, so it is held at low impact and is reversible with the politics. Source: Policy_Regulatory.

IMPACT 0.60

27 Rate Transmission Is Inverted

TAILWIND

Market

NVIDIA holds \$41,657M of net cash at 5.4% gearing with long-term debt down 32% over four years, so higher policy rates raise its interest income rather than its costs and it carries no meaningful refinancing exposure. The conventional corporate rate transmission runs backwards for this balance sheet. Source: Rates_and_Credit.

IMPACT 0.65

WAVES Secular trends (multi-year)

19 factors

WAVE + (8)

1 Reasoning-Era Inference Demand

WAVE +

Internal

Inference is shifting from one-shot prediction to reasoning and agentic workloads, which consume orders of magnitude more compute per query and scale with USAGE rather than with capex approvals. This is the only demand driver that grows without a customer budget decision, and it is why management describes demand as accelerating even at \$96.2B of quarterly revenue. Source: Key_Drivers

IMPACT  0.95

2 Annual Architecture Cadence

WAVE +

Internal

Five named generations committed through 2030 — Blackwell, Rubin (full production since 1 June 2026), Rubin Ultra H2 2027, Feynman 2028, Rosa Feynman 2029-2030 — plus a second annual cadence for the Groq-derived LPU line. A custom-ASIC programme taking two to three years to field aims at a target that moves twice before it arrives. AUDIT: reclassified from tailwind to wave_positive — five committe

IMPACT  0.85

3 Networking Attach Compounding

WAVE +

Internal

Data Center networking grew 199% year over year to \$14.8B in Q1 FY2027 against compute's 77%, and FY2026 networking exceeded \$31B, more than 10x fiscal 2021 when Mellanox was acquired. Networking scales faster than GPU count as coherent clusters grow, and it is the layer with the fewest credible substitutes. Source: Segment_Revenue_Mix

IMPACT  0.85

4 Founder-Led Capital Discipline

WAVE +

Internal

ROIC of 68.8% trailing, diluted share count down 4.2% since FY2022 despite \$7.2B of stock compensation, buybacks suspended during the pending ARM deal then resumed at 274% of free cash flow through the FY2023 downturn, and CEO pay falling 27% in a year the stock rose 24%. Compensation is keyed to non-GAAP operating income, so management is not paid on investment marks. AUDIT: reclassified from tai

IMPACT  0.75

11 AI Infrastructure Buildout Wave

WAVE +

Sector

Data-centre semiconductor revenue is forecast to rise from \$477.1B in 2026 to \$843.2B by 2030 — a 15.3% compound rate, multiples of GDP growth, and enough that NVIDIA's revenue can grow without gaining a single share point. Source: TAM_and_Growth.

IMPACT  0.95

17 Sovereign AI Demand Channel

WAVE +

Sector

Sovereign AI tripled to over \$30B in FY2026, a policy-funded buyer class that lacks the volume predictability and engineering depth to justify in-house silicon and is therefore structurally resistant to the ASIC substitution eroding hyperscaler demand. Source: TAM_and_Growth.

IMPACT  0.70

19 Supply-Constrained Accelerator Pricing

WAVE +

Sector

Accelerator pricing has held at levels only a supply-constrained market permits: an estimated 83% gross margin at a \$40,000 B200 price against 64-68% for AMD's MI300X and MI355X, with NVIDIA raising prices rather than discounting into the cost shock. Source: Pricing_Analysis.

IMPACT  0.75

28 Conventional Macro Exposures Immaterial

WAVE +

Market

Measured rather than assumed: NVIDIA's own 10-K puts a 10% adverse dollar move at \$180M of pre-tax AOCI impact, 0.08% of FY2026 revenue, because it bills in dollars regardless of customer domicile. Geographic revenue concentration in this name is a policy exposure, not a currency one. AUDIT: 'commodity exposure' here means CLASSICAL TRADED commodities (energy, metals), which a fabless model does n

IMPACT  0.60

● WAVE - (10)

5 Inference Share Loss to Custom Silicon

WAVE -

Internal

Custom ASIC shipments growing 44.6% against 16.1% for merchant GPUs, with parity projected by 2027, concentrated in exactly the stable-model inference workload a purpose-built part wins. Hyperscale — the cohort building TPU, Trainium, Maia and MTIA — is 50.9% of company revenue. This is the single structural threat every recent large bet is defending against. AUDIT: reclassified from existential_t

IMPACT  0.90

6 Earnings Quality Drift

WAVE -

Internal

Cash conversion fell to 69.7% and FCF conversion to 65.8%, the weakest in five years, while DSO rose to 76 days and the accruals ratio doubled to 18.3%. Roughly two thirds of the deterioration is unrealized investment gains inflating the net-income denominator rather than operational decay — adjusted, conversion returns to about 83.6%. The genuine operational concern is receivables growing 63.9% a

IMPACT  0.75

7 Opex Leverage Exhausted

WAVE -

Internal

Of the 27.9-point EBIT margin gain since FY2022, 18.1 points came from operating expenses falling from 27.6% to 9.5% of revenue. With SG&A already at 2.1% and R&D at 7.8%, that lever is largely spent — EBIT margin should now track gross margin closely rather than expanding independently, which also means it reverses harder on a revenue stall. AUDIT: reclassified from headwind to wave_negative. Exh

IMPACT  0.70

8 Key-Man Concentration

WAVE -

Internal

No succession plan exists and the board has not required one at a company of roughly \$5.5 trillion market value. Roughly sixty executives report directly to the CEO. Continuity risk concentrates in the Rubin roadmap, the CUDA developer relationship and the hyperscaler relationships held personally. Low probability in any year, severe and abrupt if triggered. AUDIT: reclassified from headwind to wa

IMPACT  0.60

12 Training-to-Inference Workload Shift

WAVE -

Sector

Inference is projected to reach two-thirds of all AI compute spending by 2026 and 70-80% by 2028-2030, and inference is the workload class whose price sensitivity and volume predictability justify a custom ASIC. The mix shift expands units while eroding the share that carries the most margin. Source: Industry_Trends.

IMPACT  0.90

13 Custom ASIC Substitution

WAVE -

Sector

Custom ASIC shipments grew 44.6% year over year against 16.1% for merchant GPUs, ASIC-based servers are on track for 27.8% of 2026 accelerator server shipments, and Broadcom alone carries a \$73B AI backlog against a \$100B FY2027 target. AUDIT: this is the INDUSTRY-SIDE lens on the same force the internal circle carries as Inference Share Loss to Custom Silicon; the two are paired and are now label

IMPACT  0.90

14 Memory and Foundry Cost Inflation

WAVE -

Sector

Memory pricing has turned extreme and TSMC has pushed double-digit increases, and pass-through is only partial industry-wide — NVIDIA has guided its own gross margin to a 71-72% trough before settling at 72-73% in fiscal 2028. AUDIT: reclassified from headwind to wave_negative — management guides the margin concession out to fiscal 2028, which is a multi-year path rather than a transient cost spike

IMPACT  0.80**16 Hyperscaler Capex Turns Debt-Funded**

WAVE -

Sector

Meta priced \$30B of debt plus a ~\$27B off-balance-sheet SPV, Alphabet raised ~\$25B then ~\$31B more, and Amazon's trailing-4Q capex now exceeds its trailing-4Q operating cash flow at 102% — the demand signal for the whole supplier chain has acquired a financing dependency, with depreciation drag flagged as the 2027-2028 risk. AUDIT: demoted 0.85 to 0.80. This force is carried twice — here as the in

IMPACT  0.80**22 Hyperscaler Capex Financing Shift**

WAVE -

Market

NVIDIA's revenue is its customers' capital expenditure, and that capex has moved into the bond market — about \$125.9B of on-balance-sheet issuance by Meta, Alphabet and Amazon in under five months, plus a roughly \$27B off-balance-sheet SPV, with Amazon's trailing-4Q capex at 102% of its trailing-4Q operating cash flow. Demand has acquired a credit dependency it did not have two years ago. Source:

IMPACT  0.90**25 Two Diverging Chip Markets**

WAVE -

Market

China's semiconductor localisation has shifted from long-term ambition to urgent industrial mobilisation, producing not one global AI chip market but two diverging ones. This bounds NVIDIA's addressable market permanently rather than probabilistically, and unlike the Strait risk it is already happening. Source: Geopolitical.

IMPACT  0.70● **EXISTENTIAL THREAT** (1)**21 Taiwan Strait Concentration**

EXISTENTIAL THREAT

Market

NVIDIA's own 10-K discloses that manufacture of its product components and final assembly are concentrated in Taiwan and Korea among other locations, and that the ultimate impact of that consolidation is unknown — a primary-source concentration risk with no second source before 2027. The widely quoted sizing of that risk (a six-month TSMC outage at roughly \$75B of earnings) rests on ONE anonymous

IMPACT  0.90

MUD Ambiguous / unclear

5 factors

MUD (5)

9 Supply Ceiling on Volume

MUD

Internal

Advanced wafers at TSMC and HBM across SK Hynix, Micron and Samsung are the two binding constraints named by Investor Relations, and the FY2028 growth guide of about 70% is explicitly supply-constrained against customer forecasts implying a doubling. Scarcity currently protects margin through allocation rather than discounting; it reverses into a pricing test whenever it clears. AUDIT: demoted fro

IMPACT 0.70

10 Self-Funded Demand Loop

MUD

Internal

About \$60B of cash deployed into equity investments since FY2026 (\$17,502M FY2026 plus \$42,400M H1 FY2027) and roughly \$39B of deal consideration across Groq, Hugging Face and Poolside — all directed at NVIDIA's own customer base. Returns arrive as non-operating income rather than revenue. Not circular revenue in the accounting sense, but a genuine correlation problem: the portfolio amplifies the

IMPACT 0.90

15 Export Control Regime

MUD

Sector

BIS moved H200-class chips off blanket presumption of denial to case-by-case review effective 15 January 2026, yet licences granted from February 2026 have produced no revenue to date and Chinese import permission remains unresolved. Direction is genuinely unknowable today. AUDIT: this is the INDUSTRY-WIDE lens — the January 2026 rule change names AMD's MI325X alongside NVIDIA's H200, so controls

IMPACT 0.70

18 Sector Multiples Ignore Business Quality

MUD

Sector

Regressing forward P/E on quality score across this peer group yields R-squared of 0.007 excluding Micron on n=4, i.e. NO usable relationship between business quality and multiple — and two market-data sweeps on the same date put the peer median at 21.9x or 30.8x, so whether NVIDIA sits at parity or at a 29% discount is not determinable. Direction unknowable today, which is why this is mud rather

IMPACT 0.70

24 US Export Control Regime

MUD

Market

The realised damage is taken and bounded — China fell to 9.1% of FY2026 revenue and a \$4.5B H20 charge is behind the company — but the forward direction is genuinely unknowable: BIS moved to case-by-case review in January 2026 while a new rule targeting Chinese access to remote AI servers is reported. Classified mud to match the sector circle and this bag's own undetermined fork. AUDIT: this is th

IMPACT 0.75

Contents

NVDA — Investment Recommendation


What the Company Does	13
Business Model	16
Key Financials	19
Competitive Position	22
Growth and Risks	25
Key Insights	28
Lexo Analysis	30
Trading Analysis	41

APPENDIX

Appendix I	48
Appendix II	63
Appendix III	68
Frequently Asked Questions	81
Mind Map Summary	85
Glossary	87
Lexo Terminology	91

What the Company Does

NVIDIA stopped selling a component into someone else's system and started selling the system — which is why one platform is now 92.5% of the company, and why the interesting question is no longer what it sells but who it sells to



The unit of sale is a co-designed cabinet, not a part. Revenue mix is Q2 FY2027, the quarter ended 2026-07-26.

The developer ring is the asset that carries no balance-sheet value and does not transfer in an acquisition.

The single most useful fact about what NVIDIA does today is the unit of sale. A decade ago it sold a chip that a server maker designed around. Today Data Center products ship as rack-scale systems in which the GPU, the CPU, the NVLink switch, the DPU, InfiniBand and Ethernet networking, the thermal envelope and the CUDA software stack are specified as one object. Supply-chain intelligence puts a Vera Rubin NVL72 cabinet at \$5–7 million per unit and the Rubin Ultra NVL144 at \$7–8.8 million, neither confirmed by NVIDIA. That framing is the only one in which the reported economics are coherent: the customer is buying a delivered cost per token, so the price of the silicon inside the cabinet is not the decision variable. It is also why gross margin ran at 74.7% on a trailing-twelve-month basis to Q2 FY2027 against a peer median of 58.4%, and EBIT margin at 65.2% against 21.5% — a hardware business does not earn those numbers selling parts.

THE SCALE, AND HOW FAST IT ARRIVED

Revenue was \$60,922M in FY2026's predecessor year FY2024, \$130,497M in FY2025, \$215,938M in FY2026 (ended 2026-01-25) and \$302,969M on a trailing-twelve-month basis to Q2 FY2027. That is 8.0x in four fiscal years and a further 83.4% on the trailing measure. ROIC was 68.8% TTM. The company did not build this capability in the window in which it was paid for it: CUDA shipped in 2006, the AlexNet result that validated it landed in 2012, and the revenue that justified both arrived after late 2022. For sixteen years NVIDIA carried a general-purpose compute platform on the P&L of a gaming-card business, and cumulative R&D since inception is \$76.7B against TTM R&D of \$23,592M — 7.8% of TTM revenue. No competitor was willing to carry that cost for that long, which is the origin of the position rather than a consequence of it.

WHAT IS ACTUALLY SOLD, IN THE DISCLOSURE THAT EXISTS TODAY

Data Center — \$89.0B in Q2 FY2027, 92.5% of the \$96,221M quarter, up 117% year over year. There is no sense in which this is a diversified business; every other line is a rounding error against this one. From Q1 FY2027 NVIDIA splits it into **Hyperscale — \$49B, 50.9% of company revenue** — and **ACIE (AI clouds, industrial and enterprise) — roughly \$40B, 41.6%**, derived as Data Center less Hyperscale. The split exists to show that the faster half of the platform is no longer the handful of hyperscalers who are simultaneously funding competing silicon.

Networking, the line inside Data Center that compounds fastest. Data Center networking was \$14,800M in Q1 FY2027, up 199% year over year against 92% for the segment as a whole; full-year FY2026 networking exceeded \$31B, more than 10x fiscal 2021, the year Mellanox was acquired. Networking is what makes 72 GPUs behave as one machine — the piece a customer cannot drop without surrendering the system performance it bought.

Edge Computing — roughly \$7,221M in Q2 FY2027, 7.5% of revenue, derived as total less Data Center. From Q1 FY2027 this bucket carries GeForce RTX gaming GPUs alongside automotive, robotics, workstations and AI-RAN. Gaming — the business that built the company, \$16,000M and 7.4% of revenue in FY2026 — is no longer separately disclosed.

Software, which is small in revenue and decisive in retention. CUDA is given away and monetised indirectly; enterprise software is sold through NVIDIA AI Enterprise. It is not a material revenue line, and reading it as one misreads it. Its job is to make the hardware non-substitutable.

A naming caution that matters for any period comparison: NVIDIA describes itself three ways at once — two GAAP operating segments (Compute & Networking, Graphics), the market platforms above, and the 10-K end markets. The Q1 FY2027 re-segmentation means pre- and post-change series are not directly comparable, and **no TTM segment column exists at all** because the platform disclosure was discontinued. This memo carries that gap rather than bridging it; the fullest available series is FY2022 to FY2026, in Appendix: Historical_Financials.

WHO BUYS — THE CONCENTRATION THAT SETS THE CEILING ON THE MULTIPLE

For fiscal 2026, one direct customer was **22% of total revenue** and a second **14%**. In the first half of FY2027 three direct customers were **16%, 15% and 13%** — roughly 44% of revenue across three counterparties, each of which also funds an internal alternative to NVIDIA's product. The improvement is real and modest. ACIE is management's answer to it, and the answer is only as good as ACIE's independence: several of its largest names — CoreWeave, Nebius — are companies NVIDIA has itself invested in, and the frontier labs in that bucket are financed by the same hyperscalers sitting in the other one. The diversification is genuine at the invoice level and thinner at the level of who is ultimately funding the purchase. That thread is picked up in the next section, because it has stopped being a disclosure footnote and become part of how the business model works.

THE CONSTRAINT, AND WHERE THE ECOSYSTEM SITS

Investor Relations has named the two binding constraints in as many words: advanced wafers, which depend primarily on TSMC, and high-bandwidth memory from Micron, SK Hynix and Samsung, with CoWoS packaging at TSMC the recurring practical bottleneck. Neither is diversifiable on a short horizon. While supply binds, NVIDIA allocates rather than discounts — which is why margin is expanding into a memory cost shock rather than compressing. The cost of holding that position is visible: inventory stood at \$31,575M at 2026-07-26, and the 10-K states plainly that NVIDIA places non-cancellable orders and pays deposits to secure capacity. The template for what a demand pause does to this P&L already exists — the April 2025 H2O licence change produced a **\$4.5B charge** for excess inventory and purchase obligations in Q1 FY2026, a write-down of committed supply rather than a slow fade in revenue.


Against that sits the asset with no carrying value: more than 4 million developers, more than 40,000 companies

using NVIDIA AI technologies and 15,000 startups in the Inception programme. It does not appear in the accounts and does not transfer in an acquisition, and it is the thing most exposed to a portable compilation layer that made CUDA optional.

The bridge to the next section. Everything above describes a company that has made its product harder to unbundle and its customer base narrower at the same time. What it does not yet explain is the newest part of the model — that in the last twelve months NVIDIA began putting its own balance sheet behind the customers placing the orders. That is where the business model, rather than the product, has actually changed.

Business Model

NVIDIA now runs two business models stacked on each other — a rack-scale product model that earns 65.2% operating margins, and a twelve-month-old balance-sheet model that funds part of its own customer base and pays out as non-operating income



Loop 1 is the business everyone underwrites. Loop 2 is about a year old, correlates with Loop 1 rather than hedging it, and its scale is not established — the earlier combined figure was withdrawn.

LOOP 1 — HOW THE PRODUCT MODEL ACTUALLY EARNS

The product model has three revenue layers and one pricing mechanism. **Rack-scale systems** are where the money is: GPU, CPU, NVLink switches, DPUs, network interface cards, InfiniBand, Ethernet and software architected together so that system performance, power efficiency and cost per token improve at the same time. **Networking** is the fastest-growing layer and the least appreciated — \$14,800M in Q1 FY2027, up 199% year over year. **Software** is immaterial as revenue and decisive as retention.

The pricing mechanism is the annual architecture cadence. Committing to large annual gains in performance per watt and performance per dollar lets NVIDIA charge a premium for the current generation while obsoleting the prior one on a published schedule. It cuts both ways, and this is the underappreciated half: Rubin promises up to a **10x reduction in inference token cost versus Blackwell**, which is a 10x deflation applied to NVIDIA's own installed base. Unit growth has to outrun that permanently.

The second unusual property is that **supply, not demand, sets the revenue line**. At the Q2 FY2027 print on 2026-08-26 management guided FY2028 revenue growth of approximately **70%** and stated explicitly that the outlook is supply-constrained, against customer forecasts implying more. While that holds, gross margin is protected — 75.0% GAAP in Q2 FY2027 against 71.1% for FY2026 — backlog is deep, and reported growth understates demand. It also means the first quarter in which supply catches up is the first quarter the pricing

test actually happens.

The third is that the customer base is being deliberately re-weighted, for the reason the previous section set out: roughly 44% of revenue sits with three direct counterparties (16%, 15% and 13% in the first half of FY2027), each funding an internal alternative. ACIE at roughly \$40B in Q2 FY2027 is the hedge.

WHERE THE CASH GOES — AND THE POINT AT WHICH THE MODEL CHANGES SHAPE

Cash from operations was **\$134,360M TTM** against capex of **\$7,354M** (2.4% of TTM revenue), leaving free cash flow of **\$127,006M**. On a trailing basis NVIDIA returned **\$62,091M** to shareholders — \$55,315M of buybacks and \$6,776M of dividends after the quarterly rate went from \$0.01 to \$0.25 — and spent **\$23,592M** on R&D. In June 2026 it also issued seven series of senior notes maturing 2028 to 2056, lifting total debt to **\$33,366M** against **\$228,984M** of equity at 2026-07-26. A company generating \$134,360M of operating cash flow does not need external funding for operations.

The fourth use of capital is the one that changes the description of the business. **Equity investment has overtaken R&D as a use of cash inside twelve months:** \$17,502M of equity purchases in FY2026 and **\$42,400M in the first half of FY2027 alone**, roughly \$60B of cash deployed, plus roughly **\$39B of deal consideration** committed across Groq (about \$20B, structured as a non-exclusive technology licence), Hugging Face (reported \$11.9–12.9B) and Poolside (\$6B licence plus \$1B equity) — against a prior decade in which NVIDIA's largest completed acquisition was Mellanox at about \$7B. The holdings are its own customers: OpenAI, CoreWeave, Nebius, Nokia.

FUNDING ITS OWN DEMAND IS NOW PART OF THE BUSINESS MODEL — STATED AT THE SCALE THE EVIDENCE SUPPORTS

This is the sharper of the two findings in this section, and it is easy to overclaim, so state the mechanism and the limit separately.

The mechanism, which the workbook holds as established. Capital goes out to customers as equity; product goes out to the same customers as revenue; receivables come back; and the equity stakes mark up because those customers raised at higher valuations — partly on the strength of the compute they bought. The portfolio therefore **amplifies the operating cycle instead of diversifying it**. Four numbers locate it in the accounts: non-operating income of **\$32,164M TTM** against \$11,063M for all of FY2026, roughly **17% of TTM net income** and mostly unrealised marks; ACIE revenue of roughly \$40B in Q2 FY2027, the bucket holding the AI clouds NVIDIA has invested in; receivables of **\$63,059M**, up 63.9% against revenue growth of 40.3%; and cash conversion of **69.7% TTM**, the weakest in five years against 85.6% in FY2026. That is why this memo tells the reader to underwrite the 65.2% operating margin and not the 63.7% net margin: the operating line is clean, the net line is partly a function of private-company marks.

What this is not. It is not circular revenue in the accounting sense. NVIDIA is not booking revenue on sales funded by its own cash in a way that nets to zero; the customers are real, the deployments are real, and the receivables are being collected. The disclosure is complete and the compensation plan deliberately excludes these gains.

The limit, which the workbook enforces. The *magnitude* claim has been withdrawn. An earlier version of this workbook aggregated the positions into a single figure of roughly **\$140B**; that number added a mark-to-market valuation (\$99B of portfolio value, itself containing \$23.7B of largely unrealised first-half gains) to cash outlays, and it is retracted. Cash deployed (~\$60B of equity purchases plus ~\$39B of deal consideration) and portfolio value (\$99B) are kept on two separate bases and are never summed. The \$500B-plus third-party financing platform with Apollo, BlackRock, Blackstone, Brookfield, Goldman Sachs and KKR is a **capital-mobilisation target, not committed or deployed capital, and none of it is NVIDIA's own** — it is excluded from every total here. So: the correlation mechanism survives the workbook's own stress test; **the scale at which it moves the revenue line is not established**, and the probe that would have established it failed. The steelman rates the

rebuttal to this argument **Plausible, not Strong**, and the factor is carried on the ecosystem circle as *mud* — present and directionally unknowable — at a weight of **7 of 100**, the second-smallest of the eight promoted factors. That weighting is the honest expression of a real mechanism whose size nobody in this corpus can measure.

What would resolve it. Two observable things. First, whether the third-party financing platforms grow faster than NVIDIA's own equity book — third-party capital funding customer capacity is genuine market expansion; NVIDIA's own capital funding it is a concentrated bet dressed as growth. Second, and cleaner, a quarter in which **operating income rises while net income falls**. That inversion cannot happen while the marks are appreciating, and it will happen in exactly the conditions that pressure the core business.

THE STRATEGY UNDERNEATH THE SPENDING

Read the recent deals as a list and they look unfocused — an inference-chip licence, an interconnect programme for rival silicon, an open-model repository, a hardware-abstraction startup. Read them against the threat model and they are one campaign, defending one flank: stable-model inference at scale, which is where a purpose-built part beats a general-purpose one, and where custom ASIC shipments have been growing at roughly **44.6% year over year against 16.1% for merchant GPUs**. Groq buys a different inference architecture rather than a better GPU. **NVLink Fusion** monetises the silicon NVIDIA loses — if a hyperscaler's XPU plugs into NVLink, NVIDIA earns on the rack regardless of whose accelerator sits in it, with the AWS collaboration the first named instance. Hugging Face owns the point where open models are distributed. OctoAI, at \$250M the smallest of the four, is the most revealing: a portable compilation layer is the one thing that would make CUDA optional, and NVIDIA bought one rather than let it mature elsewhere. The cost of that decisiveness is concentration — about \$39B of deal consideration in twelve months, none of it yet carrying outcome evidence.

The bridge to the next section. Both loops end in the same place: the financial statements. The next section shows what they have actually produced — and why two of the numbers in it move in opposite directions.

Key Financials

Metric	FY2024	FY2025	FY2026	TTM Q2 FY2027	Note
SCALE AND PROFITABILITY — memo summary; the full FY2022–TTM series with 18 rows, segment and geography is in Appendix: Historical_Financials					
Revenue (\$M)	60,922	130,497	215,938	302,969	8.0x in four fiscal years, then +83.4% again on the trailing measure.
Revenue growth YoY (%)	126.0%	114.0%	65.0%	83.0%	FY2023's +0.2% (not shown here; see appendix) is the only observed downside case.
Gross margin (%)	73.0%	75.0%	71.0%	75.0%	FY2026's 71.1% is the Blackwell ramp; Q2 FY2027 alone was 75.0% GAAP.
EBIT margin (%)	54.0%	62.0%	60.0%	65.0%	The cleanest profitability read — unaffected by investment marks.
Net margin (%)	49.0%	56.0%	56.0%	64.0%	TTM 63.7% is inflated by \$32,164M of non-operating income. Underwrite the 65.2%, not this.
Diluted EPS (\$/share)	1.19	2.94	4.9	7.91	FY2022 in the appendix is restated for the June 2024 10-for-1 split.
Free cash flow (\$M)	27,021	60,853	96,676	127,006	CFO less capex in every column.
FCF margin (%)	44.0%	47.0%	45.0%	42.0%	Peaked at 46.6% in FY2025. THE DIVERGENCE: this fell 2.9 points over the last year while EBIT margin rose 4.8.
ROIC (%)	64.0%	89.0%	71.0%	69.0%	NOPAT / (total debt + equity – cash), measured at period end. 68.8% TTM.
FINDING 1 — OPERATING LEVERAGE, NOT PRICING, DROVE THE MARGIN					


METRIC	FY2024	FY2025	FY2026	TTM Q2 FY2027	NOTE
R&D (% of revenue)	14.0%	10.0%	9.0%	8.0%	R&D quadrupled in dollars (\$5,268M FY2022 to \$23,592M TTM) while falling from 19.6% of revenue to 7.8%.
SG&A (% of revenue)	4.0%	3.0%	2.0%	2.0%	8.0% of revenue in FY2022 to 1.7% TTM — selling a \$5–7M rack costs little more than selling a \$30,000 server.
Opex intensity, R&D + SG&A (% of revenue)	19.0%	13.0%	11.0%	9.0%	Equals the gross-margin-to-EBIT-margin spread by construction: 27.6 points in FY2022, 9.5 points TTM.
EBIT margin gain, FY2022 to TTM (pts)				27.9	37.3% to 65.2%.
of which gross margin (pts)				9.8	About 35% of the gain — this is the pricing-and-mix half.
of which lower opex intensity (pts)				18.1	About 65% of the gain. SG&A cannot fall much below 2%, so this source is largely spent; EBIT now tracks gross margin near point-for-point in both directions.
FINDING 2 — WORKING CAPITAL NOW CONSUMES THE CASH FLOW					
Cash from operations (\$M)	28,090	64,089	102,718	134,360	14.8x over the period.
Changes in working capital & other (\$M)	-6,727	-15,392	-26,578	-69,447	DERIVED (CFO less net income, D&A and SBC). TTM is 2.6x the FY2026 drag and 23% of TTM revenue; it also absorbs ~\$23.7B of unrealised investment gains, so it overstates the true working-capital effect.
Accounts receivable (\$M, period end)		23,065	38,466	63,059	2.7x in eighteen months against revenue that roughly 2.3x'd. The 10-Q attributes it to extended payment terms on large multi-quarter agreements.

METRIC	FY2024	FY2025	FY2026	TTM Q2 FY2027	NOTE
Inventory (\$M, period end)		10,080	21,403	31,575	3.1x — pre-buying scarce HBM and CoWoS capacity under non-cancellable commitments.
DSO (days sales outstanding)	59.9	64.5	65	76	Stable near 60–65 days for four years, then 76.0.
Cash conversion, CFO / net income (%)	94.0%	88.0%	86.0%	70.0%	69.7% TTM is the weakest in the series; about 83.6% with the investment marks stripped out.
Accruals ratio, (NI – CFO) / assets (%)	3.0%	8.0%	8.0%	18.0%	Doubled to 18.3% in two quarters.
QUALITY OF THE REPORTED NET LINE					
Non-operating income (\$M)			11,063	32,164	\$15,929M of the TTM figure landed in Q1 FY2027 alone — largely mark-ups on the equity portfolio, not operations.
Non-operating income / net income (%)			9.0%	17.0%	Roughly 17% of TTM net income is non-operating. From Q1 FY2027 non-GAAP no longer excludes stock-based compensation (\$6,386M in FY2026), removing the last easy adjustment.
BENCHMARK — TTM vs peer median (Appendix: Peer_Comps)					
	NVIDIA TTM	Peer median	Gap (pts)		
Gross margin (%)	75.0%	58.0%	16.3		Percentages in columns B and C; column D is a difference in percentage points.

Competitive Position

The share NVIDIA lost is the share it chose not to defend — and the workbook still refuses to call that proven

Between 2023 and 2026 NVIDIA gave up roughly **11-12 percentage points** of data-centre AI accelerator revenue share, from about 92% in 2023 to about 81% in 2026. Over the aligned window its gross margin went **up**, from 72.7% in FY2024 to 74.7% on a trailing-twelve-month basis to Q2 FY2027. Companies losing a competitive fight discount; this one raised prices into a memory cost shock. That pair of facts is the whole competitive argument — and the limit of the argument is that it is still only *consistent* with the bullish reading, not decisive for it.



The rack is the unit of sale. NVIDIA is conceding the accelerator die at the top of the stack and defending the fabric and software beneath it. The bottom bar is the offset that has to hold for the trade to be a trade.

WHERE THE SHARE ACTUALLY WENT — NOT TO THE NAMED RIVAL

The standard framing of NVIDIA versus AMD and Intel measures the wrong contest. AMD's entire data-centre segment recorded **\$5.8B in Q1 calendar 2026**, up 57% year over year, against NVIDIA's **\$75.2B** of data-centre revenue in the quarter ended 26 April 2026, up 92%; the absolute gap widened by roughly **\$34B in that single quarter**. AMD's accelerator share sits at **5-7%** as of 2026. Intel's data-centre revenue of **\$5.1B** in the comparable quarter is mostly CPUs, and NVIDIA's roughly 4% equity holding in Intel makes that relationship closer to partnership than rivalry.

The share went instead to custom silicon designed for NVIDIA's own largest customers. Custom ASIC shipments grew **44.6% year over year** against **16.1%** for merchant GPUs (TrendForce, May 2026), and ASIC-based servers are on track for **27.8% of accelerator server shipments in calendar 2026**. Broadcom, which co-designs

Google's TPU and Meta's MTIA, carries a **\$73B cumulative AI backlog** against a **\$100B FY2027 AI revenue target** — roughly three times AMD's annualised data-centre run-rate of about \$23B, and the correct comparison once the backlog is stated on a like basis. This is the single largest external factor in the composition, carried at **weight 20 of 100** as Share Ceded to Custom Silicon.

THE OFFSET NVIDIA IS BEING PAID, AND WHAT IT IS WORTH

NVIDIA's response is not to contest the die. It is to defend the rack and the fabric: an equity stake in Marvell in March 2026 with Marvell supplying custom XPU's compatible with NVLink Fusion, **\$3.5B into MediaTek** on the same logic, and NVLink Fusion opened so that a third-party accelerator still lands inside an NVIDIA cabinet. The economic evidence is the networking line — data-centre networking of **\$14.8B in Q1 FY2027, up 199% year over year**, against **92%** growth for the data-centre segment as a whole, compounding at roughly twice the rate of the business it attaches to. Sized against the ceded accelerator revenue, that is approximately **\$39B a year of incremental networking against about \$26B a year foregone**. The same lock-in shows up in content pricing: platform revenue opportunity per gigawatt has moved **\$18B (Hopper) to \$25B (Blackwell) to \$40B (Vera Rubin)**, and compute fell from 91.5% to 82.4% of data-centre revenue over four quarters as the rest of the cabinet took the difference. That is the ecosystem factor Rack-Scale Platform Lock-In, **weight 17 of 100**.

WHY THIS MEMO HOLDS THE READING AT PLAUSIBLE, NOT STRONG

The two facts above do not settle the question, and the workbook's own steelman declines to upgrade them. **Networking is not gated by high-bandwidth memory supply; compute is.** So a quarter in which HBM allocation caps accelerator units while the fabric ships freely produces exactly the observed pattern — networking outgrowing compute, share slipping, margin holding — under *either* reading. The evidence does not discriminate between "NVIDIA is deliberately trading a low-margin socket for a higher-value attach" and "NVIDIA is supply-capped on the part that matters and the fabric is filling the gap." Three further limits are on the record: the fabric strategy is roughly two quarters old; the 199% growth rate is measured off a small base; and no disclosure in this corpus states what a rack containing a third-party accelerator earns NVIDIA relative to one containing its own. The margin evidence is also already softening — management guided gross margin down to **72-73% for fiscal 2028** from 75% in the quarter just reported, a roughly three-point structural give-back. The falsifier is specific and observable each quarter: **two consecutive quarters in which data-centre networking growth comes within roughly 10 percentage points of overall data-centre segment growth while accelerator share keeps falling**. That pattern breaks the offset arithmetic directly. The kill switch above it is a frontier training run of consequence completing on non-NVIDIA silicon at acceptable utilisation, which this corpus does not show happening at frontier scale.

WHAT THE PEER TABLE SAYS ABOUT PRICE VERSUS QUALITY

On operating quality NVIDIA is not close to its peer group: TTM gross margin **74.7% against a 58.4% peer median**, EBIT margin **65.2% against 21.5%**. On price it is a modest premium, and on two distinct measures that must not be collapsed into one. Forward **P/E of 24.7x against a 21.9x peer P/E median** is a **12.8% premium**. Forward **EV/EBITDA of 20.3x against a 17.7x peer EV/EBITDA median** is **14.7% above**. Those are different rows with different denominators. Both are also true alongside the fact that 20.3x forward EV/EBITDA sits **below NVIDIA's own five-year low of 40.0x** against a five-year average of 45.9x — cheap against its own history, a modest premium against its peers. For the full seven-name table and the TSMC memo row, **see Appendix: Peer_Comps**.

The peer comparison has a hard boundary worth naming rather than filling: no peer ROIC, net margin or EV/Revenue figure is obtainable from this corpus, and no TTM growth figure for Marvell. Those rows are absent from the appendix because they could not be sourced, not because they were unfavourable.

WHAT THIS LEAVES FOR THE SECTIONS THAT FOLLOW

The competitive position is a **mix** question, not a defeat question, and the mix question is unresolved on the evidence available as of 8 September 2026. It is also not the question that sets the price. The next section takes the growth and risk picture, and the section after that localises the valuation gap to the single forecast row that carries it.

Growth and Risks

Growth is no longer the question — whether customers can be **funded** is, and the risk that decides the downside is deliberately not in the model

For the first time in NVIDIA's history the near-term constraint is supply, not demand: management guided FY2028 revenue growth of approximately **70%** at the Q2 FY2027 print on 26 August 2026, and said customer forecasts point higher still, with the gap being what it can build. That inverts how the drivers below should be ranked. They are ordered by durability, not by size today.

THE FOUR DRIVERS THAT STILL COMPOUND ONCE SUPPLY CLEARS

1. The workload shift from one-shot inference to reasoning and agentic AI. Inference is projected at **two-thirds of all AI compute spending by 2026** and **70-80% by 2028-2030**. A reasoning model consumes far more compute per query than a single forward pass, and an agent issues many queries per task — which converts inference from a cost to be optimised down into a volume that scales with usage. It is the only driver that grows without anyone approving a new capital budget.

2. Committed platform revenue against a quantified market. Management has stated visibility to **\$0.5 trillion of Blackwell and Rubin revenue through calendar 2026**. Third-party corroboration runs the same way: cloud industry backlog above **\$2 trillion**, top-five hyperscaler capex expected near **\$800B in calendar 2026** and **\$1.3 trillion in 2027**, and data-centre semiconductor revenue forecast at **\$477.1B in 2026** rising to **\$843.2B by 2030** — a **15.3% compound annual rate**. At that market growth NVIDIA needs no share gain at all to grow.

3. The annual cadence, and Vera Rubin. Vera Rubin ships in volume in the second half with roughly **10x lower inference token cost than Blackwell**, and the platform's revenue opportunity per gigawatt moves from **\$25B (Blackwell) to \$40B (Vera Rubin)**. A custom ASIC programme that takes two to three years to field is aiming at a target that has moved twice by the time it arrives.

4. Customer broadening. Hyperscalers fell to roughly **half of data-centre revenue in Q2 FY2027**, with AI clouds, industrial, enterprise and sovereign buyers carrying the rest; sovereign AI alone **tripled to over \$30B in FY2026**. That class has neither the engineering depth nor the volume predictability to justify custom silicon, so it is structurally NVIDIA-shaped demand. It is also the one driver an investor can verify independently, by watching whether it outgrows hyperscale for two consecutive quarters.

MACRO RISK HERE IS CONCENTRATED, NOT DIFFUSE

This is a measurement finding, not a stylistic preference, and it is the reason this section carries no generic macro paragraph. **Conventional macro exposures are immaterial by measurement.** A 10% adverse dollar move produces a **\$180M** pre-tax impact in accumulated other comprehensive income — **0.08% of FY2026 revenue of \$215,938M**. Rates are not merely small but *inverted*: NVIDIA holds **\$41,657M of net cash** against total debt of **\$8,468M** at **5.4% gearing**, so higher rates raise interest income rather than costs. Commodity exposure is structurally absent because NVIDIA is fabless. Bolting a standard macro sensitivity layer onto this model adds noise and, on the rate term, the wrong sign.

What replaces those variables is three channels a conventional screen has no column for. **Customer capital budgets:** two direct customers were **22% and 14% of FY2026 revenue**, and three were **16%, 15% and 13% in H1 FY2027**. **One island:** treated separately below. **Upstream input cost:** memory and foundry pricing transmit straight into gross margin, and management's guided give-back from 75% to a **72-73% FY2028** settling point is worth roughly **\$9.1B a year** at about \$3.0B of gross profit per margin point — roughly fifty times the FX exposure. None of the three is hedgeable, and that is what distinguishes them from the ones that are.

THE EXTERNAL RISKS, AND THE 54% OF THE CASE OUTSIDE NVIDIA'S CONTROL

Three factors sit on the external circle of the composition and together carry **54 of the 100 weight points** — the share of this case that no amount of management execution touches. The largest is **Price, Terminal Value and the Guidance Bet (weight 21)**, argued in the valuation sections. The second is **Share Ceded to Custom Silicon (weight 20)**, argued above. The third is the one that changed in kind during 2026: **Demand Is Credit-Funded (weight 13)**.

The mechanism is a substitution, not a slowdown. The four US hyperscalers bought **\$433.9B of property and equipment in the four quarters through March 2026** against roughly **\$149B of reported depreciation** — and then issued debt anyway: Meta **\$30B in October 2025** plus a roughly **\$27B off-balance-sheet SPV**, Alphabet about **\$25B in November 2025** and about **\$31B more in February 2026**, Amazon **\$24.9B** after **\$15B** — roughly **\$150B in under five months**, all while generating record operating cash flow. **Amazon's trailing-four-quarter capex now equals 102% of its trailing-four-quarter operating cash flow**. Above 100%, every incremental dollar of build is externally funded by definition. NVIDIA has extended credit support of its own: an obligation capped at **\$105 billion in aggregate** on roughly 4.25 gigawatts of IT load under 20-year leases. Revenue a supplier helps finance is not the same asset as revenue a customer funds from profits — and the enforcement point is a tenor mismatch, with GPU-backed facilities written against customer contracts averaging roughly three years while the debt and leases run five to twenty. The named durability risk is the **2027-2028 depreciation drag**.

THE TWO INTERNAL RISKS THAT ARE ALREADY IN THE NUMBERS

Operating Leverage Already Spent (weight 5, intrinsic circle). Of the **27.9 percentage points** of EBIT margin gained between FY2022 and FY2026, **18.1 points** came from operating expense falling from **27.6% to 9.5% of revenue**, with SG&A from **8.0% to 2.1%**. That lever is spent, and it runs backwards at the same gearing: FY2023 is the observed proof — revenue of **\$26.97B, up 0.2%**, while GAAP diluted EPS fell **55%** and EBIT margin went from **37.3% to 15.7%**.

Founder-Led Execution, Unplanned Succession (weight 8, intrinsic circle). The premium here is attached to a person and is repriced in one step rather than accrued back gradually. Over the decade to 2026, still-founder-led companies returned about **15.4% a year** against **6.9%** for companies whose founder-CEO had left. NVIDIA's 10-K carries only the generic succession clause — that failure to ensure effective succession planning "could hinder our strategic planning, execution, and long-term success" — which is disclosure of the risk, not disclosure of a plan. The tradable part of this theme is the difference between those two, and NVIDIA sits on the less-informative side of it.

WHY THE TAIWAN STRAIT IS DELIBERATELY HELD OUTSIDE THE FACTOR COMPOSITION

It is the only exposure in this workbook that is existential, and it is **not** one of the eight promoted factors. That is a decision, not an oversight.

The exposure is real and unmitigated: over **90% of the world's most advanced chips** are made in Taiwan, over **90% of TSMC's advanced wafer capacity** sits there, no viable second source exists for 3nm or 2nm at scale before **2027**, TSMC's Arizona capacity does not reach mass production until **2028**, and advanced packaging stays in Taiwan longer still. A six-month TSMC disruption is estimated at a **\$75B earnings impact — 38.9% of TTM net income** — though every quantified figure in that estimate rests on a single anonymous source and should be treated as indicative; the concentration itself is 10-K disclosed.

The reason it stays outside the composition is that **it does not discriminate between NVIDIA and any alternative use of the same capital in this complex**. It is shared identically by AMD, Broadcom, the hyperscalers and every custom-silicon programme; no company-level analysis changes it and no company-level action reduces it; and it is binary and unobservable, with a limited blockade able to halt exports overnight

without military engagement and therefore without warning. A factor set exists to explain why *this* security is priced where it is relative to others. An exposure that moves every name in the complex by the same amount explains nothing about relative price. **It is a position-sizing question, not a stock-selection one** — it belongs in how much of the portfolio is committed to the AI complex at all, and it is carried in this workbook as the only existential threat on the lightcone and as a named kill switch, not as a weight in the factor composition. Folding it in would have produced a factor that is large, unfalsifiable and identical for every alternative — which is the definition of a weight that does no work.

Key Insights

1. VARIANT VIEW — we are not betting against the market's enthusiasm. We are betting against the company's own guidance.

Seven findings carry this case into the Lexo analysis. The first is the variant view, and it is not the one a reader expects from a **Hold** on a stock that has re-rated.

The **-40.6%** gap between the probability-weighted intrinsic value of **\$136.78** and the current price of **\$230.36** does not diffuse across the model. It localises to one row: **FY2028 revenue of \$474,927M, a 26.8% growth rate**, against consensus of **\$681,500M** and management's guided **~70%**, issued at the Q2 FY2027 print on 26 August 2026. Consensus and management are close to each other; this workbook is the outlier. That reframes the entire position. The conventional bear case is that the market is exuberant. Ours is narrower and more falsifiable: **management's FY2028 number is too high, and consensus is repeating it**. The same disagreement is already visible one year earlier — FY2027 consensus of **\$401,270M** against this workbook's **\$374,401M** is only **-6.7%** at the headline, but H1 actuals of **\$177,800M** dilute it; on the unreported half the gap is **-13.8%**, and the implied Q3 FY2027 of roughly **\$97,600M** sits **9.6% below the company's own \$108.0B guide**.

2. THE DISAGREEMENT HAS A DATE ON IT: 2026-11-17.

This is the rarest property in the case and the reason conviction is **Low** rather than negative. A thesis that resolves on an announced date is not a view that has to be held indefinitely. On **17 November 2026**, at the Q3 FY2027 print, management either reiterates FY2028 near **70%** with backlog detail — in which case the model row is wrong and the position re-opens — or it does not. **FY2028 guidance below 45%** is the short/exit trigger; **two of the four big hyperscalers guiding capex flat or down in the late-January-2027 window** is the corroborating one. Between now and then the recommendation is **no new weight initiated**, with a **\$195.00** stop on two consecutive weekly closes, **-15.3%** from \$230.36.

3. THE MARKET HAS REPRICED THE GROWTH RATE AND ALMOST NONE OF THE UNCERTAINTY.

The de-rating has already happened: forward **EV/EBITDA of 20.3x** sits **below NVIDIA's own five-year low of 40.0x** against a five-year average of **45.9x**. But the market's pricing of *dispersion* has not moved with it — implied volatility sits at an **IV rank of 0-9 out of 100** and short interest is **1.23% of float**. The multiple says the market lowered its growth expectation; the options and the borrow say almost nobody thinks the outcome is in doubt. Reverse-solved, **\$230.36 requires 6.95% perpetual growth, or a 44.2% five-year revenue CAGR**. That is the asymmetry, and it is in the uncertainty rather than in the direction.

4. THE STRONGEST BULL ARGUMENT IN THIS WORKBOOK IS HELD AT PLAUSIBLE, NOT STRONG — BY THE WORKBOOK'S OWN STEELMAN.

NVIDIA ceded roughly **11-12 percentage points** of accelerator share since 2023 while gross margin rose **72.7% to 74.7%**, and networking compounds at **199%** against **92%** for the data-centre segment. The mix reading — that this is a deliberate trade, not a defeat — is the load-bearing argument for the quality verdict, and it is **not upgraded**, because **networking is not gated by high-bandwidth memory supply while compute is**, so the observed pattern is produced equally by a supply cap on the part that matters. The offset arithmetic (**~\$39B/yr incremental networking against ~\$26B/yr foregone**) supports the reading without settling it, because the gross margins of the two streams are not separately disclosed. A memo that reported this as settled would be reporting something the research refused to conclude.

5. THE PEER GROUP PAYS NOTHING FOR QUALITY, SO THE RISK SITS IN THE EARNINGS RATHER THAN IN THE RATING.

NVIDIA ranks **2nd of 7** on gross, EBIT and FCF margin and **2nd of 6** on growth, while ranking **4th of 7** on forward multiple. Regressing forward multiple on quality across that peer set produces an **R-squared of 0.007** — no relationship at all. The practical consequence is that the usual de-rating bear case has little to work with: at a **12.8% P/E premium** and a **14.7% EV/EBITDA premium** to the peer medians, there is not much multiple to give back. What can move is the **E**, which is exactly where finding 1 points. For the full table, see Appendix:

Peer_Comps.

6. EARNINGS QUALITY IS DRIFTING, AND IT IS THE STRONGEST INTERNAL ARGUMENT AGAINST THE QUALITY VERDICT.

Cash conversion fell to **69.7%** and FCF conversion to **65.8%**, the weakest in five years, with DSO at **76 days** and the accruals ratio doubled to **18.3%**. About two-thirds of that drop is a denominator effect from **\$32.2B of non-operating investment gains** — adjusted, conversion is roughly **83.6%** — but the receivables and inventory build is real, and the investment book that produces the gains grew from roughly **\$7B to \$99B in a year**. Those gains correlate with the demand cycle rather than hedging it: a quarter in which operating income rises while net income falls is the signal the portfolio has turned, and it will turn in the same conditions that hurt the core business.

7. THE BINDING CONSTRAINT IS SUPPLY, NOT DEMAND — WHICH IS A QUALITY SIGNAL AND A BLIND SPOT AT THE SAME TIME.

The limiter is qualified, shippable high-bandwidth memory rather than theoretical capacity. A company whose ceiling is what it can build rather than what it can sell has a demand problem no competitor would trade for. But a supply-constrained guide also means **NVIDIA is not currently being tested on whether it can win orders at the margin** — it cannot fill the orders it has. The first quarter in which management stops describing the outlook as supply-constrained is the first real test of pricing power, and it precedes any deterioration in the reported numbers. Management has already published the ceiling it expects to concede: gross margin settling at **72-73% in fiscal 2028**, down from **75%** in the quarter just reported.

WHAT THESE SEVEN SET UP

This is a position question, not a quality question. 54 of the 100 factor weight points sit outside NVIDIA's control, the Taiwan Strait exposure is held outside the composition entirely as a position-sizing matter, and the one row that carries the **-40.6%** gap is a forecast this workbook built **price-blind and guidance-blind** — it never saw the 70% guide, and the appendix names FY2028 as the most likely error in its own model. The Lexo analysis that follows grades the business on its merits (**Diamond in Air | Ocean | Luminous**, N-of-1 **83/100**, multibagger gate **BORDERLINE**) and the trading analysis prices the position (**Hold**, conviction **Low**, **Bull 25% / Base 50% / Bear 25%**). Both are answering the same question from opposite ends: an exceptional business at a price that requires management to be right.

Lexo Analysis

What the business is made of

FOUNDATION ASSESSMENT

The quality question is closed. NVIDIA scores **Diamond, 72/100** on the Lexo intrinsic axis, **Luminous, 72/100** on value extraction, **Ocean, 70/100** on market position and **N-of-1, 83/100** on tier; only the extrinsic environment (**Air, 56/100**) is short of its ceiling. Headline verdict: **Diamond in Air | Ocean | Luminous**. The consequence matters more than the labels: if quality is closed, then *every live argument about NVIDIA is an argument about something else*. Diligence spent re-testing unit economics is spent in the wrong place. What remains open is position — the pod that follows — and price, settled against the model in **Appendix: DCF_Summary** and against a date, the Q3 FY2027 print on **2026-11-17**.

Three anchors, all TTM to Q2 FY2027 (quarter ended 2026-07-26). **Return on invested capital of 68.8%** on gearing of 5.4%, so the spread over any defensible cost of capital holds whatever the cost of equity proves to be — and sustained, against a five-year average of 53.4%. **A gross margin of 74.7% against a 58.4% peer median** and an **EBIT margin of 65.2% against a 21.5% peer median**, roughly three times the median and the widest quality gap in this workbook; the peer set is the six companies in NVIDIA's FY2026 10-K competitor group, and EBIT is the honest line because, unlike net margin, it is untouched by the \$32,200M of trailing non-operating investment income. And **compound revenue growth of 100.0% over FY2023–FY2026**, \$26,974M to \$215,938M, carried to a TTM revenue base of \$302,969M.

The qualitative test agrees: NVIDIA raised prices into a memory cost shock rather than discounting, and no challenger has replicated the stack — AMD took the die, Broadcom the custom design, Groq the inference runtime; one layer each. Foundation Quality subtotals **27 of 30** on the N-of-1 scorecard, with ROIC Quality the rubric's only **10 of 10**.

THE FORK THAT WAS WEIGHED: DIAMOND HELD, GOLD CONSIDERED

This is the first year in five with evidence behind a downgrade, so the fork belongs in the record. Three facts argued Gold. **Cash conversion fell to 69.7% TTM**, the weakest of the five-year series (85.6% at FY2026); FCF conversion fell to 65.8%; the accruals ratio doubled from 8.4% at FY2026 to 18.3% TTM; DSO extended from 65.0 to 76.0 days against four prior years inside a 51.8–65.0 day band. **Operating leverage is spent** — of the 27.9 points of EBIT margin gained since FY2022, 18.1 came from opex falling from 27.6% to 9.5% of revenue, and at SG&A of 2.1% nothing further is available. And **management has published a ceiling**, guiding gross margin of 72–73% for fiscal 2028 against 75% most recently reported — the first quantified limit on pass-through it has conceded.

The counter-argument that won is that none of the three touches the moat. Roughly two-thirds of the conversion fall is a denominator effect — strip the \$32,200M of non-operating income out of net income and cash conversion is about 83.6%, inside its historical range. Capex is not the cause at 2.4% of revenue, days inventory of 38.0 sits well below the FY2023 spike of 69.8, and deferred revenue rose 149% in two quarters to \$6,412M — customers prepaying. A Diamond becomes Gold when its switching cost stops rising, not when its margin falls three points. **Diamond is held at 72/100 — the low end of the 67–100 band, and low deliberately.** What would move it: cash conversion below 70% for two further quarters, or the guided 72–73% floor breaking.

VALUE EXTRACTION: LUMINOUS TODAY, AND WHAT WAS RULED OUT

Extraction asks how much of that value reaches shareholders; most of it does. **A free-cash-flow margin of 41.9% TTM against a 22.0% peer median. \$62,091M returned in twelve months** — \$55,315M of buybacks plus \$6,776M of dividends, 48.9% of free cash flow — with diluted shares down 4.2% since FY2022 to 24,285M TTM despite \$7,241M of stock-based compensation. Pricing on value delivered is still rising: revenue opportunity per gigawatt moved from roughly \$18B for Hopper to \$40B for Vera Rubin, a compound 49.1% per generation.

A step-down to **Radiant** was weighed, on the theory that the \$99B equity portfolio (marked at 2026-07-26) was corrupting the extraction reading. It was investigated and **found not attributable to the investment book**. ASU 2016-01 routes unrealised equity marks through the income statement, so the book inflates the *denominator* of every net-income ratio — \$23,707M of gains in H1 FY2027 alone — but it does not touch free cash flow, the capital return or the FCF margin — the extraction measures. The genuinely operational residue is DSO at 76.0 days: one metric, not a class change. **Luminous, 72/100, ships today**. The modal three-year case still steps this axis to Radiant — a deliberate sequence, not a hedge: monetisation is optimised, conversion is deteriorating, and they are different properties.

SO WHAT

An intrinsic Diamond at 72 and a Luminous 72 mean the business cannot be the disagreement. Nothing in the foundation explains why a probability-weighted intrinsic value of **\$136.78** sits **40.6% below** the **\$230.36** price. The gap is not in the margins, the returns or the capital discipline; it is in what this business is assumed to be *positioned* to sell three years out, and in one FY2028 revenue row (**Appendix: DCF_Summary**). Foundation hands the first of those forward: what position does this business hold, and is the platform what keeps it?


The number that looks bad

PLATFORM ASSESSMENT

The rack-scale platform is the only reason this position reads Ocean while unit share falls. It carries the ecosystem circle — 17 of the 100 factor-weight points, the largest inside NVIDIA's own control — and it has to be argued, because the raw share numbers read the other way.

NVIDIA's platform: five layers, and the one competitors attack

Platform Strength 23 of 30 — Lexo classification, September 2026



Every credible challenger attacks layer 4 alone. NVLink Fusion now admits third-party silicon into layer 3.

Switching cost rises with height: NVIDIA owns the top three layers, contests the fourth, does not own the fifth.

Share of AI accelerator revenue fell from roughly 92% in 2023 to **81.1% in 2026** — about **10.9 percentage points** — and structurally, not cyclically: custom ASIC server shipments grew **44.6% year over year** against **16.1%** for merchant GPUs. NVIDIA already supplies about **69% of the \$281B core intelligent-data-centre segment**, so growth cannot come from share. **Market Position Durability scores 6 of 10**, the N-of-1 rubric's weakest criterion and the reason Platform Strength subtotals **23 of 30**. And it went in-house rather than to a merchant rival (AMD holds 5–7%, its whole data-centre segment \$5.8B against NVIDIA's \$75.2B in the same quarter): **in-house substitution is arguably worse** — an internalised function removes the willingness to pay permanently, where a rival can be out-executed.

WHY THE AXIS STILL READS OCEAN, 70/100

Because the unit of sale changed underneath the share statistic. In 2023 a buyer bought accelerators and wired them together; in 2026 it buys a coherent cabinet — a Vera Rubin NVL72 quoted to OEMs at **\$5M–\$7M per unit**, a Rubin Ultra NVL144 at **\$7M–\$8.8M**, bundling 72 GPUs at 288 GB of HBM4 each with NVLink switching, SuperNICs, DPUs and liquid cooling into one warranted machine. The switching cost is the rack, not the chip: packages drawing toward **1,800W** force 800V DC distribution, so changing vendor means re-engineering power, cooling and floor layout, not swapping a card — on top of roughly **4 million CUDA developers** and about **60% of TSMC CoWoS capacity**.

The offset that makes Ocean arithmetic rather than assertion sits in the mix. Compute fell from **91.5% to 82.4% of Data Center revenue in four quarters** while networking grew **199% TTM against 92% for the segment** (Q4 FY2026: +263% against +58%). That is roughly **\$39B a year of incremental networking revenue against about \$26B a year of accelerator revenue foregone at the old share** — and NVLink Fusion admits third-party XPU's into NVIDIA racks, converting a rival's die into channel. The market is still filling, at **15.3% a year to \$843.2B by 2030**.

The fork, and how far it was taken. The temptation is to grade that rebuttal Strong; the workbook's own steelman refuses. The offset is revenue, not margin — no margin comparison between ceded accelerator revenue and gained networking revenue exists in this corpus — and networking is not HBM-gated while compute is, so the evidence cannot discriminate between "NVIDIA replaced the revenue" and "NVIDIA sold what it could supply." Held at **Plausible, not Strong**. Not River, though: River is timing-led, which was NVIDIA in 2019, not a company at 81.1% share with supplier allocation priority.

THE MEDIUM IT OPERATES IN: AIR, 56/100

The one axis short of its ceiling, and the one execution cannot fix. Tailwinds are real — sovereign AI tripled past **\$30B** of policy-funded demand, and **\$41,657M of net cash** inverts rate transmission so higher rates raise interest income. It falls short of Space on friction outside management's control: Taiwan concentration disclosed in NVIDIA's own 10-K, with no viable 3nm or 2nm second source before 2027; export controls that took China from **13.1% to 9.1% of revenue** with a **\$4.5B** charge taken; and memory, foundry and packaging costs running roughly **\$9.1B a year** of gross profit at TTM revenue.

SO WHAT

Ocean at 70 inside Air at 56, intact only because the platform moved the unit of sale up a level faster than share was lost at the chip. A four-point fall in the total, which Market Position Durability could deliver alone, takes the tier from **83 to 79** and out of the N-of-1 band. **What would move it: a frontier training run of consequence completing on non-NVIDIA silicon** — the contested layer reaching the layers above it. Whether the stack is

irreproducible or merely hard to reproduce is the next pod's question; the price of getting it wrong sits in

Appendix: DCF_Summary.

The prosecution brief: what earns the 83

N OF 1 ASSESSMENT

The tier is earned and it is not insulated. NVIDIA scores **83 of 100** on the N-of-1 rubric, clearing the 80 threshold by **three points** — and the multibagger gate, which asks a different question, is held at **BORDERLINE**. The two verdicts are not in tension. The 83 says nobody has rebuilt what this company holds. The borderline says nobody can yet show that what it holds compounds from here.

Foundation Quality, 27 of 30. The rubric's only perfect mark is **ROIC Quality, 10 of 10 — 68.8%** trailing return on invested capital against gearing of **5.4%**, sustained against a five-year average of **53.4%**. **Moat**

Irreproducibility takes 9 of 10, held below full marks only because a portable compilation layer would attack it directly. **Pricing Power takes 8 of 10** for a specific and recent reason: management has published its own ceiling, guiding gross margin of **72–73% for fiscal 2028** against **75%** most recently reported — the first quantified limit on pass-through NVIDIA has ever conceded.

Platform Strength, 23 of 30 — the weakest category and the binding constraint on the score. **Platform Layer, 9 of 10** (CUDA, NIMs, NeMo and Omniverse above the silicon; NVLink, Spectrum-X and full racks below it). **Orbital Potential, 8 of 10**, held there because the **\$99B** ecosystem portfolio is an asset and a risk at once. And **Market Position Durability, 6 of 10** — the lowest criterion anywhere in the rubric, on roughly **11 points** of accelerator share ceded since 2023.

Transcendence Factors, 33 of 40. Category-Defining Potential, 9 of 10: NVIDIA created the category and the programming model it runs on, and per-gigawatt monetisation moving **\$18B** (Hopper) → **\$25B** (Blackwell) → **\$40B** (Vera Rubin) says the category is still being redefined upward. **8 of 10** each on historical precedent, competitor failure history and network-effect compounding — nothing earns a 10, the honest read on growth unprecedented in scale inside an industry that has produced violent up-cycles before.

WHO TRIED, AND WHICH LAYER EACH ONE STOPPED AT

AMD took the die — 5–7% of accelerator share, a whole data-centre segment of **\$5.8B** against NVIDIA's **\$75.2B** in the same quarter, matching on silicon and losing on software utilisation (roughly **45%** model FLOPs utilisation against **50–55%**). **Intel took the accelerator** — a decade of Gaudi, still single-digit training share. **Broadcom took the custom design** — a **\$73B** AI backlog against a **\$100B** FY2027 target, the most serious of the four and still one layer. **Groq took the inference runtime.** Four credible challengers, four single layers, no stack — and NVIDIA itself could not buy the layer it lacked, its **\$40B** ARM bid blocked and Grace built internally instead.

THE DEFENCE BRIEF, AND THE FORK THAT WAS WEIGHED

Three objections, none dismissed. **Share is falling for a third consecutive year** — custom ASIC server shipments grew **44.6%** year over year against **16.1%** for merchant GPUs. **The network effect is uneven** — four million CUDA developers is a real moat in training and a demonstrably weaker one in inference, which is where roughly **two-thirds** of AI compute spending is going by 2026. **And the environment is doing some of the work** — a market compounding at **15.3% a year** flatters any leader, and this company's own **FY2023** (revenue **+0.2%**, EBIT margin **37.3% → 15.7%**, ROIC to **14.9%**) shows the business without the wave.

A downgrade to **NFU-Diamond-in-Rough (60–79)** was genuinely on the table, because all three objections are true. It was declined because none of them touches irreproducibility, which is the only property the tier measures: a company can lose share and remain N-of-1 if what it holds cannot be rebuilt, and here the switching cost is still rising while unit share falls. The margin must be carried honestly, though — **83 minus 4 is 79**, so a **four-point** deterioration drops the band, meaning Market Position Durability falling 6 → 2 or two-point falls in two criteria. An earlier draft claimed a two-point fall sufficed; 83 minus 2 is 81, still inside. Carry the margin, not just the label.

WHY THE MULTIBAGGER GATE STAYS AT BORDERLINE

Three axes at their top class — **Diamond, Ocean, Luminous** — inside a **15.3%** compounding market is the profile a multibagger track normally requires. Two things hold it short of YES.

The modal path steps an axis down rather than up. The **50%** base case moves Value Extraction from **Luminous to Radiant** as cash conversion settles in the high sixties to low seventies, and reaching the five-year bull state requires inference share loss to *stabilise*, which no evidence in this workbook yet shows.

The runway is constrained by penetration, not market size. NVIDIA already supplies roughly **69% of the \$281B** core intelligent-data-centre segment — about **31 points** of headroom — and about **40.6% of the wider \$477.1B** data-centre semiconductor market, about **59 points** on that denominator. The **15.3%** rate applies to the *wider* market, so the two must never be read together. A compounding outcome depends on the addressable definition widening, not on taking the rest of a segment NVIDIA already dominates. **What resolves it:** networking attach growth continuing to outpace segment growth while share stabilises → **YES**; networking converging on segment growth for two consecutive quarters while share keeps falling → **NO**.

CALIBRATION, IN ONE PARAGRAPH

Cisco, 1997–2001 is the most-cited parallel and the most over-applied: it bites only through the financing channel, and Cisco's customers were leveraged carriers with no cash flow where NVIDIA's four largest are among the most creditworthy issuers in the market. **Intel, 1995–2015** is the instructive one — Intel's moat failed not to a competitor but when the ecosystem migrated to a different compute model, precisely the portable compilation layer this workbook names as its highest-severity threat. **ASML** is the bull case's end state, not its current state, and **Micron** is the live warning: highest quality score in this peer set, bottom of its rating.

SO WHAT

The rarity is real and measurable at 83 of 100; the runway depends on which market you count. Neither is the reason NVIDIA trades where it does. The tier holds, the gate is borderline, and the **40.6%** intrinsic gap sits in none of it — it sits in one revenue row three fiscal years out. That is the next pod, and it is the argument this memo exists to make.

The variant view, stated plainly and uncomfortably

INVESTMENT CASE

NVIDIA is an exceptional business with a position question, not a quality question — and a price question that resolves on a known date, 2026-11-17. Everything above this pod closed the quality question: a **68.8%** ROIC, a **74.7%** gross margin against a **58.4%** peer median, three Lexo axes at their ceiling, a tier of **83 of 100**. None of it is disputed anywhere in this workbook. What is disputed is one revenue row three fiscal years out — and the identity of the counterparty on the other side of it.

The 40.6% gap between this workbook's value and NVIDIA's price is a bet against management's own FY2028 guidance — not against market exuberance.

The row is this. Our model carries **FY2028 revenue of \$474,927M, growth of 26.8%**. Consensus carries **\$681,500M**. Management, at the Q2 FY2027 print on **2026-08-26**, guided to roughly **70%**. We are **43 points** below the company's own number and **30.3%** below the street's.

That is an unusual and uncomfortable place to stand, and the memo should say so rather than dress it up. The comfortable variant view on a stock like this is that the crowd is euphoric and the multiple absurd. Neither is available. Forward EV/EBITDA is **20.3x** — **below NVIDIA's own five-year low of 40.0x** and far under its **45.9x** five-year average; forward P/E of **24.7x** is a **12.8% premium** to the **21.9x** peer P/E median, and forward EV/EBITDA sits **14.7% above** the **17.7x** peer EV/EBITDA median. Cheap against its own history, a modest premium against peers. The market has already conceded most of the deceleration argument; what remains is a disagreement with the issuer.

THE FORK THAT WAS WEIGHED HERE — AND IT IS THE MEMO'S MOST CONSEQUENTIAL ONE

The projections model in this workbook was built **price-blind and guidance-blind**, by architecture rather than by accident: the platform forbids the projections bag from reading a market price, because knowing the price while setting assumptions biases them toward it. The consequence is that the model never saw the **~70%** guide. It was built from this workbook's own sector evidence — accelerator share drifting into custom inference silicon against a total market compounding at **15.3% a year** to 2030.

Two options were live once the guide appeared. **Option one: revise FY2028 toward guidance.** That closes the gap almost entirely without touching another input, and it is what a reader who trusts the company would do.

Option two: leave the model and report the disagreement in full. The counter-argument decided it — retrofitting a price-blind model to a figure it was built without destroys the only independent estimate in the workbook, and a valuation that agrees with guidance by construction tells a reader nothing the press release does not. Option two was taken, at a real cost: **this workbook's own view is that FY2028 is the most likely error in the model**, and it is publishing a valuation it believes contains that error. Management calls the guide supply-constrained rather than demand-constrained, with **\$500B+** of disclosed visibility behind it, and its guidance hit rate in this corpus is **3 of 3** — clean, on a thin sample.

The FY2027 agreement is thinner than its headline too: the **-6.7%** gap is diluted by **\$177,800M** of H1 actuals that cannot differ, and on the unreported half it is **-13.8%** — implying roughly **\$97,600M** of Q3 FY2027 revenue against the company's **\$108.0B** guide, **9.6% below**.

UNDERNEATH THE PRICE DISAGREEMENT, A BUSINESS DISAGREEMENT HELD AT PLAUSIBLE

The second variant view is narrower: the share decline is read as competitive defeat when the evidence supports a mix decision. Share fell roughly **11 points** since 2023 while gross margin **rose from 72.7% to 74.7%**; it went to custom silicon inside NVIDIA's own customers, not to AMD (whose whole data-centre segment did **\$5.8B** against NVIDIA's **\$75.2B** in the same quarter); and networking compounds at **199%** against **92%** for the segment — roughly **\$39B a year** of incremental networking revenue against about **\$26B a year** of accelerator revenue foregone at the old share. This workbook holds that reading at **Plausible, not Strong**, for the reason the next pod develops.

THE NUMBERS, AND THE FACT THAT NONE OF THEM CLEARS THE PRICE

SCENARIO	PROBABILITY	INTRINSIC VALUE	VS \$230.36
Bull	25%	\$189.80	-17.6%
Base	50%	\$139.24	-39.6%
Bear	25%	\$78.83	-65.8%
Probability-weighted	100%	\$136.78	-40.6%

The probabilities sum to exactly **100%** and are identical to those carried in Scenario_Analysis; the **25%** bear weight is a **15%** competitive route and a **10%** geopolitical one. **Even the bull case sits 17.6% below the current price of \$230.36.** That, not the bear case, is the finding. A three-scenario set in which the best modelled outcome still loses money is either a valuation that is right or a model whose growth assumption is wrong — and this memo has already named which of the two it suspects.

Two disciplines on those figures. **-40.6% is the intrinsic-value gap** — what the business is worth on full convergence, which is not a twelve-month assumption — while the **-13.1% probability-weighted twelve-month return** in Scenario_Analysis is what the price is likely to do inside the tactical horizon. Different measures on different clocks, never interchangeable. And the model's weak point is disclosed rather than smoothed: **71.6% of base-case enterprise value is terminal value**, breaching the 40–70% band this platform sets. Reverse-solved, today's **\$230.36** requires a **6.95%** perpetual growth rate or a **44.2%** five-year revenue CAGR against **24.8%** base and **31.3%** bull — not impossible, but what a buyer at today's price is underwriting (see **Appendix: DCF_Summary** for the assumptions, the FCF bridge, the valuation walk, the reverse-solve and the scenario table).

SO WHAT

The rating is **Hold**, at **Low** conviction, with **no new weight initiated**. A Sell would assert this workbook's revenue path over the company's own guidance with no adjudicating evidence, against a business this same workbook rates N-of-1 at **83/100**. A Buy would ignore a valuation that fails in all three modelled scenarios. The memo's case is therefore narrow, and stated as such: *the FY2028 row is lower than the company says, that is the most likely thing this workbook has got wrong, and a dated event settles it*. What the strongest opponent says to that is the next pod.

The bear thesis at full strength: the moat protects the shrinking half of the market

LEXO STEELMAN

The strongest case against this memo is not that NVIDIA is overrated. It is that the single rebuttal the variant view leans on is rated Plausible, not Strong — and it is this workbook that refuses to upgrade it. Two of five bear arguments get Strong answers; three get only Plausible ones. That tally is worse than an earlier draft claimed, and the honest version is the one that belongs here.

A short-seller would not argue that NVIDIA is a bad business. They would argue its moat is real and pointed at the wrong territory. **CUDA's switching cost binds hardest on frontier training, and training is the shrinking half.** Inference is projected at roughly **two-thirds** of AI compute spending by 2026 and **70–80%** by 2028–2030 — price-sensitive, architecturally simpler, volume-predictable: the profile that justifies a custom ASIC.

The evidence is not thin: share went from about **92% in 2023 to 81.1% in 2026**, custom ASIC server shipments grew **44.6%** year over year against **16.1%** for merchant GPUs, ASIC servers reached **27.8%** of accelerator shipments in 2026, and Broadcom carries a **\$73B** AI backlog against a **\$100B** FY2027 target. The share went in-house rather than to a merchant rival, which removes a customer's willingness to pay permanently rather than merely negotiating over it. **How it unfolds:** inference volume compounds, in-house share of it compounds faster, the switching cost stops rising, and Market Position Durability — the rubric's lowest criterion at **6 of 10** — takes the classification down with it.

THE REBUTTAL, AND WHY IT IS HELD AT PLAUSIBLE

The bull answer is the "mix not defeat" reading: NVIDIA defends the rack rather than the die — a stake in Marvell, **\$3.5B** into MediaTek, NVLink Fusion opened so third-party accelerators land inside NVIDIA systems. Gross margin **rose from 72.7% to 74.7%** while share fell, prices were raised into a memory cost shock rather than cut, and networking compounds at **199%** against **92%** for the segment: roughly **\$39B a year** incremental against about **\$26B a year** foregone.

Rebuttal strength: Plausible — not Strong, and the reason is the most important sentence here: **networking is not HBM-gated while compute is, so the evidence does not discriminate** between "NVIDIA replaced the revenue" and "NVIDIA sold what it could supply." This workbook's own Key Insight 8 holds the binding constraint to be qualified, shippable HBM rather than demand — so the share NVIDIA "ceded" may be share it could not serve. Three further weaknesses hold the rating down: the fabric strategy is roughly **two quarters** old, the **199%** rate is off a small base, and nothing in this corpus establishes what a rack containing a third-party accelerator earns NVIDIA. **This is the workbook's most important open question.**

THE OTHER FOUR, RATED HONESTLY

BEAR ARGUMENT	REBUTTAL	STRENGTH
Three straight years of share loss is what losing looks like	It went to customers' own silicon, not to AMD — whose entire data-centre segment did \$5.8B against NVIDIA's \$75.2B in the same quarter — while margin rose and prices were raised. Losing companies discount.	Strong
Demand is financed: \$125.9B of hyperscaler debt in five months plus a ~ \$27B SPV, Amazon's trailing-4Q capex at 102% of operating cash flow	NVIDIA's stakes are equity funded from \$127,006M of TTM free cash flow, and the buyers are among the most creditworthy issuers in the market	Plausible — the direction of travel is not favourable
Operating leverage is spent: 18.1 of 27.9 margin points came from opex falling to 9.5% of revenue, and guidance is 72–73% by FY2028	A 72–73% trough is still roughly twenty points above AMD's 53%	Plausible — it concedes the trajectory and wins only on the level

BEAR ARGUMENT	REBUTTAL	STRENGTH
Strip the buildout wave and this is a cyclical semiconductor business: FY2023 ran revenue +0.2% , EBIT margin 37.3% → 15.7% , ROIC to 14.9%	FY2023 also shows what survived intact; cyclical affects the <i>rate</i> , not the <i>rarity</i>	Strong

A CLAIM THIS WORKBOOK WITHDREW, AND WHAT SURVIVED IT

Honest steelmanning includes retracting your own overreach. The **fc_08 self-funded-demand-loop** factor once carried an **order-of-magnitude claim** — that NVIDIA underwrites its demand far beyond the roughly **\$140B** of counted stakes and undertakings, on a **\$500B+** financing layer plus lease backstops and purchase financing. **That claim was RETRACTED after its magnitude probe failed.** The **\$500B+** is a stated *target* for mobilising **third-party** capital through six named firms — not committed, not deployed, none of it NVIDIA's own money — and the backstop, purchase-financing and Texas-lease figures disagree across the outlets carrying them.

The correlation mechanism survived; the magnitude did not. Still at full force: roughly **\$60B** of cash into equity stakes, about **\$39B** of deal consideration, a portfolio marked near **\$99B**, **\$32,200M** of trailing non-operating income, and a CDS market that has already repriced the correlation — **+14bp in a single session, to roughly 82**. The factor keeps **7 of 100** weight on that surviving mechanism and none on the withdrawn one.

KILL SWITCHES

1. **A frontier training run of consequence completes on non-NVIDIA silicon** at acceptable utilisation — continuous watch, most likely 2027–2028. Market Position leaves Ocean; the tier leaves N-of-1.
2. **Networking attach growth converges within roughly 10 points of data-centre segment growth for two consecutive quarters while share still falls** — quarterly, a 2027 signal. Turns the Plausible rebuttal above into a wrong one.
3. **Any Taiwan Strait blockade, quarantine or sustained export interruption** — unhedgeable before TSMC's Arizona capacity reaches mass production in **2028**. The environment goes to Mud and the rest is irrelevant for the duration.
4. **Cash conversion below 65% unadjusted (or 78% adjusted) for two consecutive quarters with accruals above 20%** — current readings **69.7%**, about **83.6%**, **18.3%**. The drift would be structural rather than a denominator effect.


VERDICT

Conviction: Held — at Low. The steelman does not move the rating because the rating already reflects it: **Hold**, no new weight. What it establishes is where the fragility sits — not in the quality evidence, which survives every attack above, but in one rebuttal (HBM-gating makes the mix-versus-defeat evidence non-discriminating) and one revenue row that **2026-11-17** begins to settle. The single thing this memo is most uncertain about is whether NVIDIA *chose* to cede that share or *could not serve* it.

Reading the six axes for this company

LEXO SYMBOLOGY

Three axes sit at the top of their ladder and one does not, and the one that does not is the only one NVIDIA cannot fix by executing well. The verdict is **Diamond in Air | Ocean | Luminous**, tier **N-of-1** at **83 of 100**. Read as a sentence: an irreproducible business, holding the dominant position in its market, converting that position into shareholder cash at close to the maximum rate — operating in a medium with friction it does not control.



Schematic aid only. Every score it depicts is stated in the table below; read the text, not the picture, for the numbers.

AXIS	CLASS	SCORE	DIRECTION	WHAT IT MEANS FOR NVIDIA
Intrinsic quality	Diamond	72	flat	What the business is made of
Extrinsic environment	Air	56	flat	The medium it operates in
Market position	Ocean	70	down	The position it holds
Value extraction	Luminous	72	down	How much of the value reaches shareholders
Tier	N-of-1	83	down	How rare it is
Trajectory	Stable	42	down	Where the classification is heading

Intrinsic — Diamond, 72. The Diamond band runs 67–100, so 72 sits near its floor, deliberately. A **68.8%** ROIC on **5.4%** gearing, a **74.7%** gross margin against a **58.4%** peer median and a **65.2%** EBIT margin against **21.5%** are Diamond evidence without qualification. The score is held low because cash conversion has fallen to **69.7%** unadjusted — the weakest in five years, about **83.6%** once the **\$32,200M** of non-operating investment income is stripped from the denominator — and because management has published a gross-margin ceiling of **72–73%** for fiscal 2028.

Extrinsic — Air, 56. The one axis short of its ceiling, and the reason the headline is not "Diamond in Space." Real tailwinds: a data-centre semiconductor market compounding **15.3% a year** from **\$477.1B** to **\$843.2B** by 2030, sovereign AI past **\$30B**, and **\$41,657M** of net cash that inverts rate transmission so higher rates raise interest income. Against them, friction management cannot remove: Taiwan concentration disclosed in NVIDIA's own 10-K with no viable 3nm or 2nm second source before 2027, export controls that took China from **13.1% to 9.1%** of revenue with a **\$4.5B** charge taken, and input costs running roughly **\$9.1B a year** of gross profit.

Market — Ocean, 70, and falling. NVIDIA *is* the market at its scale — roughly **81.1%** of accelerator revenue and about **69%** of the **\$281B** core segment — but the shoreline is receding: **11 points** of share ceded since 2023, Market Position Durability at **6 of 10**. Ocean rather than River because River is timing-led, which described NVIDIA in 2019, not a company at 81% share with supplier allocation priority.

Extraction — Luminous, 72, and falling. A **41.9%** free-cash-flow margin against a **22.0%** peer median, **\$62,091M** returned in twelve months (**48.9%** of free cash flow), and per-gigawatt content pricing compounding **49.1%** a generation from **\$18B** to **\$25B** to **\$40B**. Luminous means monetisation is optimised, and it is. What is deteriorating is *conversion* — a different property, and why the modal three-year case steps this axis to **Radiant**.

Tier — N-of-1, 83. Clears the 80 threshold by three points; a **four-point** fall drops it to NFU-Diamond-in-Rough. The criterion holding it down is the same one the bear case attacks.

Trajectory — Stable, 42. The score readers most often misread. Stable does not mean safe; it means the classification is not compounding upward. On the modal path three axes hold and one steps *down* — this is a business whose classification is expected to persist rather than improve, which is precisely why the multibagger gate is **BORDERLINE** rather than YES.

TODAY → THREE YEARS → FIVE YEARS

Today:	Diamond in Air	Ocean	Luminous	
3-Year Bull:	Diamond in Space	Ocean	Luminous	(HBM yields resolve, China reopens, ASIC share stalls near 27.8%)
3-Year Base:	Diamond in Air	Ocean	Radiant	(margin lands at the guided 72-73%, conversion stays high-60s)
3-Yr Bear Rt.1:	Gold in Air	Lake	Radiant	(portable compilation layer reaches training parity + capex turn)
3-Yr Bear Rt.2:	Diamond in Mud	Ocean	Luminous	(Taiwan escalates; business untouched, cannot ship)
5-Year Bull:	Diamond in Space	Ocean	Luminous	(packaging qualifies outside Taiwan; inference share stabilises)
5-Year Base:	Diamond in Air	Ocean	Radiant	(share drifts, market outgrows the drift, conversion settles low-70s)
5-Year Bear:	Gold in Water	Lake	Radiant	(substitution reaches training AND the buildout wave matures)

The axes move independently, which is the point of keeping four of them: the base case downgrades extraction while leaving the foundation untouched, and the geopolitical bear route collapses the environment while leaving the other three exactly where they are.


SO WHAT

Three axes at their ceiling, one held back by geography, and a trajectory that is Stable rather than rising: a business that is exceptional and not compounding its exceptionalism. That is a **quality** verdict and a **position** verdict, both settled. Neither is a **price** verdict — the Lexo axes are deliberately price-blind, and the 3–5 year intrinsic view stated here is not revised anywhere in the Trading Analysis that follows. What follows is the other clock: what a **6–12 month** holder should do about a business classified this way at **\$230.36** a share.

Trading Analysis

So what: the market has already repriced NVIDIA's growth rate and has repriced almost none of the uncertainty around it — the setup is **Neutral****, and the disagreement that decides it is with NVIDIA's own management rather than with market exuberance**

MARKET SETUP



Schematic aid only. Every figure it depicts is stated in the text below; read the text, not the picture, for the numbers.

THE SPRING STATE IS NEUTRAL

At **\$230.36** per share (close of 2026-09-04, re-confirmed 2026-09-08) NVIDIA sits **2.6% below** its 52-week high of \$236.54 and **40.4% above** its 52-week low of \$164.07. A stock that close to its high is not a coiled spring; a stock whose forward multiple has fallen below its own five-year trough is not an extended one either. The energy in this setup is neither spent nor stored — it is waiting for a dated event. The state is **Neutral**.

THE DISAGREEMENT LOCALISES TO ONE ROW, AND THE COUNTERPARTY IS MANAGEMENT

	THIS WORKBOOK (BASE)	CONSENSUS / GUIDANCE	GAP
FY2027 revenue	\$374,401M	\$401,270M (consensus midpoint)	−6.7%
FY2028 revenue	\$474,927M	\$681,500M (consensus midpoint)	−30.3%
FY2028 revenue growth	26.8%	~70% (company-guided)	−43 points
Forward EV/EBITDA (2026-09-08)	20.3x	45.9x five-year average	−25.6x

We do **not** disagree with the market about what NVIDIA is, and we do not disagree about what it is worth per dollar of near-term earnings. The –40.6% intrinsic gap reduces almost entirely to the FY2028 revenue row: **\$474,927M and 26.8% growth in this workbook against a consensus \$681,500M and management's own guided ~70%**, issued at the Q2 FY2027 print on 2026-08-26. The counterparty on that row is **NVIDIA's management, not market exuberance**.

FY2027 AGREEMENT IS THINNER THAN THE HEADLINE

The –6.7% FY2027 gap is diluted by H1 FY2027 actuals of **\$177,800M**, which are already reported and cannot differ. On the **unreported half the gap is –13.8%** (\$196,601M modelled against a \$228,000M consensus H2). Carried to the quarter, that implies roughly **\$97,600M of Q3 FY2027 revenue against the company's \$108.0B guide – 9.6% below**. State the caveat honestly: this model is annual and produces no quarterly path, so the implied Q3 figure is **derived from the half-year, not modelled**.

WHY OUR NUMBER IS LOWER, AND WHY IT MAY BE THE ERROR

The projections model was built **price-blind and guidance-blind**. It never saw the ~70% guide; it was built from this workbook's own sector evidence — accelerator share drifting into custom inference silicon against a total market compounding at **15.3% a year to 2030**. That is a defensible way to build a model and a dangerous way to end a debate. NVIDIA describes the FY2028 guide as supply-constrained rather than demand-constrained, with **\$500B+ of disclosed visibility** behind it. **This workbook's own view is that FY2028 is the most likely error in the model**, and that belongs on the page rather than in a footnote.

WHAT THE MARKET HAS REPRICED, AND WHAT IT HAS NOT

Forward EV/EBITDA is **20.3x** against a **45.9x five-year average** and a **40.0x five-year low** — below NVIDIA's own five-year trough. Against its peers the same date reads differently: forward P/E of **24.7x** is a **12.8% premium** to the **21.9x** peer P/E median, and forward EV/EBITDA of **20.3x** is **14.7% above** the **17.7x** peer EV/EBITDA median. Cheap against its own history, a modest premium against its peer set — both are true (see Appendix: Valuation_Multiples for the full multiple history and peer rows).

Set that against the options market. 30-day implied volatility is **31.9%** at an **IV rank of 0–9 out of 100**, and short interest is **1.23% of float**. The market has repriced the *growth rate* substantially and almost none of the *uncertainty* around the 2026-11-17 print. That gap — a de-rated multiple beside unrepriced volatility — is the tradable observation in this memo.

TWO HORIZONS, NAMED SO THEY ARE NOT MISTAKEN FOR A CONTRADICTION

- **3–5 year intrinsic view: unchanged and positive.** NVIDIA classifies as **Diamond in Air | Ocean | Luminous**, N-of-1 at **83/100**. Nothing in this section revises that.
- **6–12 month tactical stance: Hold.** A stock 2.6% below its 52-week high, whose modelled value depends on disputing a company guide that gets tested in ten weeks, is not one to add to or exit on conviction.

These are different questions on different clocks, not a contradiction. The next pod states what to do about it.

So what: the rating is **Hold**** at ****Low**** conviction and ****no position is initiated**** — because the one factor worth trading here is not direction, it is that the market has repriced NVIDIA's growth level and not the uncertainty around it, and that uncertainty is settled on a known date**

THE TRADE

THE FACTOR, STATED ONCE

NVIDIA's forward multiple has already de-rated to 20.3x EV/EBITDA — below its own 40.0x five-year low — while the implied volatility priced into its options sits at an IV rank of 0–9 out of 100, ten weeks before the single print on 2026-11-17 that adjudicates a 43-point disagreement about FY2028 revenue growth (26.8% modelled against ~70% guided).

The repricing has happened entirely in the *level* and none of it in the *variance*. Those two are normally repriced together.

Subtle. Both obvious readings fail in the same way. The bull reading — a cheap stock at 20.3x against a 45.9x five-year average — treats the de-rating as an opportunity without asking what the market learned on the way down. The bear reading — a DCF saying –40.6% — treats that gap as exuberance, which the multiple data flatly contradicts. The real observation sits between them: the market has already conceded most of the deceleration argument, and what remains in dispute is narrow, specific and dated.

Debatable. A serious investor takes the other side of either leg. Implied volatility is low *because* realised moves have compressed; the options market has overestimated NVDA's realised earnings move in **77% of the last 13 quarters**, and selling that volatility rather than buying it has paid for three years. And the FY2028 guide is described as supply-constrained with **\$500B+ of disclosed visibility** behind it — in which case there is no binary at all, only a print that confirms what the company already said.

Measurable. Every term is a sourced number: IV rank 0–9 of 100 and 30-day implied volatility of 31.9% (as of 2026-09-08); forward EV/EBITDA 20.3x against a 45.9x five-year average and a 40.0x five-year low; FY2028 growth 26.8% modelled against ~70% guided; an implied earnings move of about 5.7% for the 2026-11-17 print against a 5.1%–7.7% historical average.

Actionable inside 12 months. The resolving event is dated — **2026-11-17** — with a confirming event in the late-February 2027 reporting window. Both fall inside the horizon.

WHAT THE VALUATION ACTUALLY SAYS

Probability-weighted intrinsic value is **\$136.78** per share against a market price of **\$230.36** — **–40.6%**, with every modelled scenario, the bull included, sitting below the price (see Appendix: DCF_Summary for the assumptions, the FCF bridge, the valuation walk and the scenario table).

Keep that figure distinct from the other one this memo uses. **–40.6% is the intrinsic-value gap**, a statement about what the business is worth. The **–13.1% probability-weighted 12-month return** developed in the next pod is a statement about what the price is likely to do over the horizon. They are different measures on different clocks and are never interchangeable.

Reverse-solved, today's \$230.36 requires a **6.95% perpetual growth rate** — more than double the 3.0% assumed — or, holding terminal growth at 3.0%, a **44.2% five-year revenue CAGR** against 24.8% in our Base case and 31.3% even in our Bull. Neither is impossible. Both are what a buyer at today's price is underwriting.

THE RECOMMENDATION

Hold. Low conviction. No new weight initiated.

A **Sell** would assert this workbook's revenue path over the company's own guidance with no adjudicating evidence, against a business rated N-of-1 at 83/100. A **Buy** would ignore a defensible valuation that fails in all three modelled scenarios. Neither is supportable on what is known at 2026-09-08. The platform's own rule maps a Neutral spring to a starter position; that rule is **deliberately overridden here**, because opening new length into a distribution this workbook scores at -13.1% over 12 months would be following a rule against the analysis the rule exists to serve.

For an investor who already owns NVDA, the actionable half is different: protection covering the 2026-11-17 print and extending through the late-February 2027 window is purchasable at the cheapest implied volatility of the past year. No strikes or expiries are named here — no options chain was sourced — so the instruction is the **window** and the **volatility condition**, nothing finer.

THE COST OF BEING RIGHT ABOUT THE SETUP AND WRONG ABOUT THE TRADE

Cheap protection is not free protection. If NVIDIA prints in line, guides FY2028 near 70%, and the stock moves less than the ~5.7% implied — the outcome in 77% of the last 13 quarters — the premium is lost and holding no position means missing the move entirely. That is the price of admitting this workbook cannot adjudicate its own model's central assumption before **2026-11-17**. It is a real cost, and it is the honest one to pay at Low conviction.

So what: the bull case is what decides this section — at **\$189.80**** the modelled bull intrinsic value still sits ****17.6% below**** the ****\$230.36**** price, so this three-scenario set contains no outcome in which the business is worth what the market is paying for it**

SCENARIO ANALYSIS

THE THREE SCENARIOS, AND THE FACT THAT NONE OF THEM CLEARS THE PRICE

SCENARIO	PROBABILITY	INTRINSIC VALUE	VS \$230.36
Bear	25%	\$78.83	-65.8%
Base	50%	\$139.24	-39.6%
Bull	25%	\$189.80	-17.6%
Probability-weighted	100%	\$136.78	-40.6%

The probabilities sum to exactly **100%**, and the **25%** bear weight is itself two routes: a **15%** competitive route in which a portable compilation layer reaches frontier-training parity while the hyperscaler capex cycle turns as 2024-2025 depreciation lands in 2027-2028, and a **10%** geopolitical route in which Taiwan escalates and the business is unchanged but cannot ship. The **50%** base case is the one to read closely: the market compounds at the forecast **15.3%** a year to 2030, accelerator share drifts down a few points a year into custom inference silicon, the guided **72-73%** FY2028 gross margin lands as guided, and the company is still worth **\$139.24 - 39.6%** under the price. The weakness of this table is not the bear case. It is that the bull case loses too.

THE 12-MONTH DISTRIBUTION IS A DIFFERENT, SMALLER AND MORE SURVIVABLE NUMBER

SCENARIO	PROBABILITY	12-MONTH PRICE TARGET	RETURN	ANCHOR FOR THE TARGET
Bull	25%	\$236.54	+2.7%	Retest and clear the 52-week high on FY2028 commentary holding near 70% with backlog detail at the 2026-11-17 print
Base	50%	\$200.09	-13.1%	Drift back to the \$199-\$201 shelf where April, June and July 2026 all closed within \$1.41 of each other
Bear	25%	\$164.07	-28.8%	Full retrace to the 52-week low on FY2028 guidance below 45% , or a hyperscaler capex deceleration signal in the late-January 2027 window
Probability-weighted 12-month return	100%	-	-13.1%	-9.7% if the \$195.00 stop is honoured; best case +2.7%

Two numbers in that table coincide and are not the same measurement: the **base-case** 12-month return is **-13.1%** and the **probability-weighted** 12-month return is also **-13.1%**. The modal path and the weighted path land on the same figure because the upside tail (**+2.7%**) is too small to offset the bear leg. Read it as a coincidence of arithmetic, not as confirmation.

Keep the two headline percentages of this memo strictly apart. **-40.6% is the intrinsic-value gap** - what the business is worth against what it costs, assuming full convergence of price to value, which is not a 12-month assumption. **-13.1% is the probability-weighted 12-month return** - what the price is likely to do inside this section's tactical horizon. They are different measures on different clocks and neither substitutes for the other anywhere in this memo. Applying the **\$195.00** stop-loss trims the modelled **-13.1%** to **-9.7%**, because the bear target of **\$164.07** sits **\$30.93** below that stop and a disciplined holder never realises the full **-28.8%**.

THREE LIMITS THIS TABLE CARRIES ON ITS FACE

The bull target is capped by construction. All three 12-month targets are observed price levels from the last 52 weeks rather than derived valuations, which keeps them honest and truncates the upside at the **\$236.54** 52-week high. If FY2028 lands near the guided **~70%**, NVDA very likely re-rates above that level and this table puts no number on it. The **-13.1%** weighted return is therefore biased downward and should be read as a floor, not a point estimate.

The left tail is understated too. The geopolitical bear route carries **10%** of the weight but is not modelled as a cash-flow path, so the full **25%** bear weight sits on the competitive route's **\$78.83**. A Taiwan escalation would be more destructive than that figure, not less.

The scenario set is three, not five. No Steelman and no Market-Implied scenario exists upstream in the projections model's Valuation_Summary, so neither can be shown here. That matters because the Market-Implied case is the one that would price management's own **~70%** FY2028 guide, and its absence is why this table can state the disagreement but cannot adjudicate it. The full **WACC x terminal-growth** and **CAGR x**

margin grids behind these three values are in **Appendix: Sensitivity_Matrix**.

WHAT MOVES THE READER BETWEEN ROWS

The scenarios are separated by one dated observation. On **2026-11-17** NVIDIA reports Q3 FY2027 against a guided **\$108.0B** revenue and **74.0%** gross margin - but the print is not the catalyst, the FY2028 commentary attached to it is. Holding near **70%** with backlog detail puts the reader in the bull row; softening below **45%** puts them in the bear row; language that neither confirms nor breaks the guide leaves them in the **50%** base row, drifting toward **\$200.09**. The confirming events follow: the hyperscaler Q4 CY2026 capex guides in the **2027-01-26 to 2027-02-06** window, and the formal FY2028 guide on or around **2027-02-24**.

A distribution whose best outcome is **+2.7%** and whose weighted outcome is **-13.1%** does not support new length. The closing pod states what that means in instructions.

So what: **Hold, **Low** conviction, **no position initiated** - and because the platform's own rule would open one here, the override is the instruction that matters**

EXECUTION PLAN

FIELD	INSTRUCTION
Rating	Hold
Conviction	Low
Position action	None. No new weight initiated.
Spring state	Neutral
Platform rule for Neutral	Starter position - deliberately overridden
Stop-loss (existing holders only)	\$195.00 on two consecutive weekly closes = -15.3% from \$230.36
Re-entry trigger	FY2028 commentary holding near 70% with backlog detail on 2026-11-17
Short / exit trigger	FY2028 guidance below 45% ; or any two of the four large hyperscalers guiding capex flat or down in the late-January 2027 window
Horizon	6-12 months , tactical

WHY THE RULE IS OVERRIDDEN RATHER THAN FOLLOWED

The spring state is **Neutral**, and the platform's entry logic maps a Neutral spring to a starter position. Applied mechanically here it would open roughly a one-third weight. This memo declines it, on this bag's own arithmetic: a probability-weighted 12-month return of **-13.1%** against a ceiling of **+2.7%** in the single best of three scenarios. Opening length into that distribution would be following a rule against the analysis the rule exists to serve. **A Hold means hold. With no existing position, the correct action is no action.** The rule is a default for when exposure is wanted; it is not a reason to want exposure.

FOR A READER WHO ALREADY OWNS NVDA

This is not an instruction to sell, and the two paragraphs above do not apply to existing exposure. The discipline that does apply is a stop at **\$195.00** on **two consecutive weekly closes** - **-15.3%** from the **\$230.36** reference price. That level is set deliberately beneath the **\$199-\$201** shelf where April, June and July 2026 all closed within **\$1.41** of each other: a two-week close underneath it means the shelf has failed rather than been tested. The two-close confirmation exists so a single volatile week around an earnings print does not trigger the exit.

Protection across the same window is cheap - 30-day implied volatility sits at an **IV rank of 0-9 out of 100**, the low end of the past year. That is a reason to hedge **existing** exposure through the **2026-11-17** print and the late-February 2027 window, and it is not a reason to create new exposure. No strike, expiry or premium is named anywhere here: no options chain was sourced for this workbook, so the instruction is the window and the volatility condition and nothing finer.

THE THREE DATES THAT CHANGE THIS RATING

- **2026-11-17 (company-confirmed).** FY2028 commentary holding near **70%** with backlog detail means this workbook's central assumption is the thing that was wrong, the **20.3x** forward EV/EBITDA is genuinely low, and the rating moves to Buy at a full weight rather than a starter. This is the re-entry trigger.
- **Late January 2027 (2027-01-26 to 2027-02-06, estimated).** Any **two** of the four large hyperscalers guiding capex flat or down year over year is the most direct available threat to the FY2028 guide, because roughly half of Data Center revenue is four customers.
- **On or around 2027-02-24 (estimated).** The formal FY2028 guide. Below **45%** growth, the assumption holding up the multiple has broken and the rating moves to Sell.

Between those dates only the **\$195.00** price trigger applies, and it applies to a holder's existing position, because this memo opens none.

WHAT THIS MEMO IS SAYING, IN ONE LINE

NVIDIA is an exceptional business - **Diamond in Air | Ocean | Luminous**, N-of-1 at **83/100** on a 3-5 year intrinsic view - carrying a price question that this memo cannot settle and that **2026-11-17** can. Waiting ten weeks for a dated answer costs **+2.7%** of foregone upside in the best modelled case. Acting now, on a **-13.1%** 12-month expectation, costs more.

Appendix I

Historical Financials

METRIC	FY2022	FY2023	FY2024	FY2025	FY2026	TTM Q2 FY...	AGENT_NOT...
Revenue (\$M)	26,914	26,974	60,922	130,497	215,938	302,969	8.0x in four years, then +83.4% again on a TTM basis. The FY2023 flat spot is the crypto/gaming correction before the AI ramp.
Revenue Growth YoY (%)	61.0%	0.2%	126.0%	114.0%	65.0%	83.0%	FY2023's +0.2% is the whole bear case in one cell: this business CAN stall.
Gross Profit (\$M)	17,475	15,356	44,301	97,858	153,463	226,241	
Gross Margin (%)	65.0%	57.0%	73.0%	75.0%	71.0%	75.0%	Trough 56.9% (FY2023, on a \$2.2B inventory/supply charge); FY2026's 71.1% is Blackwell ramp cost, recovered to 75.0% in Q2 FY2027.
R&D (\$M)	5,268	7,339	8,675	12,914	18,497	23,592	Compounded at 36.9% against revenue at 68.3% — R&D fell from 19.6% of revenue to 8.6%.
SG&A (\$M)	2,166	2,440	2,654	3,491	4,579	5,070	From 8.0% of revenue in FY2022 to 2.1% in FY2026 — the signature of a very large unit of sale.
Operating Income/ EBIT (\$M)	10,041	4,224	32,972	81,453	130,387	197,579	

METRIC	FY2022	FY2023	FY2024	FY2025	FY2026	TTM Q2 FY...	AGENT_NOT...
EBIT Margin (%)	37.0%	16.0%	54.0%	62.0%	60.0%	65.0%	37.3% to 65.2% TTM. About two thirds of the 27.9-point gain came from operating leverage, only one third from gross margin.
EBITDA (\$M)	11,215	5,768	34,480	83,317	133,230	201,266	EBIT plus D&A; D&A is only 1.2% of TTM revenue, so EBITDA tracks EBIT closely. Not a figure NVIDIA reports.
Net Income (\$M)	9,752	4,368	29,760	72,880	120,067	192,879	
Net Margin (%)	36.0%	16.0%	49.0%	56.0%	56.0%	64.0%	TTM 63.7% is inflated by \$32.2B of non-operating investment gains. EBIT margin is the clean read.
Diluted EPS (\$/share)	0.39	0.17	1.19	2.94	4.9	7.91	FY2022 RESTATED for the June 2024 10-for-1 split; EDGAR still carries \$3.85 pre-split.
Cash from Operations (\$M)	9,108	5,641	28,090	64,089	102,718	134,360	
Capital Expenditures (\$M)	976	1,833	1,069	3,236	6,042	7,354	Peaked at 2.8% of revenue in FY2026. NVIDIA buys fab capacity from others, not for itself.
Free Cash Flow (\$M)	8,132	3,808	27,021	60,853	96,676	127,006	FCF = CFO less capex for every column, including TTM (\$134,360M - \$7,354M = \$127,006M).

METRIC	FY2022	FY2023	FY2024	FY2025	FY2026	TTM Q2 FY...	AGENT_NOT...
FCF Margin (%)	30.0%	14.0%	44.0%	47.0%	45.0%	42.0%	THE DIVERGENCE: EBIT margin rose 4.8 points over the last year while FCF margin FELL 2.9 points. Working capital and other absorbed -\$69,447M on a TTM basis.
ROIC (%)	28.0%	15.0%	64.0%	89.0%	71.0%	69.0%	NOPAT / (total debt + equity - cash). FY2023's 14.9% shows the cyclicalilty the current 68.8% masks.
Net Cash / (Net Debt) (\$M)	11,957	5,491	23,295	42,232	41,657	n/a	Positive = NET CASH (total debt less cash and marketable securities). Net cash every year. FY2026's \$41,657M is the figure the memo uses. TTM is n/a here: this row PREDATES the June 2026 senior-note issue that lifted total debt to \$33,366M at 2026-07-26.
SEGMENT REVENUE (\$M) — market platforms, FY2022-FY2026 ONLY							
Data Center	10,610	15,010	47,530	115,190	193,740		18.3x over four years; 39% of revenue to 89.7%.
Gaming	12,460	9,070	10,450	11,350	16,000		The business that built NVIDIA; 7.4% of FY2026 revenue. No longer separately disclosed from Q1 FY2027.

METRIC	FY2022	FY2023	FY2024	FY2025	FY2026	TTM Q2 FY...	AGENT_NOT...
Professional Visualization	2,110	1,540	1,550	1,880	3,200		Smallest legacy platform but grew 70% in FY2026.
Automotive	566	903	1,091	1,690	2,350		Compounded 42.8% and is still only 1.1% of revenue.
OEM & Other	1,162	455	306	389	619		Structurally shrinking; 0.3% of FY2026 revenue.
Data Center mix (%)	39.0%	56.0%	78.0%	88.0%	90.0%		Segment figures are published to 2-3 significant figures, so the five platforms sum to within \$29M (0.01%) of total revenue. The difference is not absorbed into any segment.
DISCLOSURE GAP — segment series ENDS at FY2026							From Q1 FY2027 NVIDIA reports only Data Center (Hyperscale and ACIE) and Edge Computing, where Edge Computing equals the sum of the former Gaming, Professional Visualization, Automotive and OEM & Other lines. FY2026 is the LAST year for which this table can be extended, so no TTM segment column exists.
GEOGRAPHIC REVENUE — FY2026 ONLY, customer-HQ basis (\$M)							

METRIC	FY2022	FY2023	FY2024	FY2025	FY2026	TTM Q2 FY...	AGENT_NOT...
United States					149,620	0.69	NVIDIA states customers headquartered outside the US were 31% of FY2026 revenue.
Taiwan					42,350	0.20	Taiwan-headquartered ODMs and system builders; NVIDIA estimates much of this is consumed by end customers in the US and Europe.
China (incl. Hong Kong)					19,680	0.09	Grew 15.0% in FY2026 through export controls. The Q1 FY2027 outlook then assumed ZERO Data Center compute revenue from China.

Peer Comps

COMPANY	TICKER	MARKET C...	REVENUE T...	REV GROW...	GROSS MA...	EBIT MAR...	AGENT_NOT...
							SUBJECT. Market cap = 24,100M diluted shares x \$230.36 (close 2026- 09-04, re- confirmed 2026-09-08) = \$5,551,676 M. Operating figures are computed from SEC EDGAR for the four quarters ended 2026- 07-26, not taken from a market-data vendor. BASIS NOTE ON THE TWO MULTIPLES: 24.7x is forward P/E on FY2027E and 20.3x is forward EV/EBITDA on FY2027E, both struck 2026-09-08 at \$230.36. The sector bag separately carries a 22.0x NTM consensus P/E struck on a different date and window; both are on the record and are NOT reconciled into one figure.
NVIDIA	NVDA	5,551.68	302.97	83.0%	75.0%	65.0%	

COMPANY	TICKER	MARKET C...	REVENUE T...	REV GROW...	GROSS MA...	EBIT MAR...	AGENT_NOT...
Broadcom	AVGO	1,700	89.1	49.0%	77.0%	49.0%	The closest true comparable and the one peer that is a structural threat rather than a merchant rival: it designs the custom XPU's that displace NVIDIA GPU's. HIGHER gross margin than NVIDIA (77.1% vs 74.7%); EBIT margin 48.6% is roughly three quarters of NVIDIA's, not half. \$73B AI backlog, \$100B FY2027 AI revenue target.
Micron	MU	1,150	90.27	167.0%	73.0%	66.0%	Commercially a SUPPLIER (HBM sold into accelerators), not an accelerator rival, but named in NVIDIA's 10-K competitor group so included on the same single criterion as every other peer. Its 167% TTM growth and 65.7% operating margin are a memory-cycle peak, not a run rate — which is why it is simultaneously the highest-quality and the cheapest name in the set.

COMPANY	TICKER	MARKET C...	REVENUE T...	REV GROW...	GROSS MA...	EBIT MAR...	AGENT_NOT...
AMD	AMD	759.83	41.31	40.0%	53.0%	16.0%	The named merchant rival. GROSS MARGIN BASIS DIFFERS: 53.2% is AMD's Q1 2026 GAAP quarter, NOT a TTM figure like NVIDIA's — this one cell is not basis-consistent with the column. EBIT margin 15.7% against NVIDIA's 65.2%; it competes by discounting 30-40% on inference, and still trades at a HIGHER forward multiple than NVIDIA.
Intel	INTC	503.38	57.03	7.0%	39.0%	8.0%	Not a factor at the high end: single-digit AI training accelerator share, 38.6% gross margin, 7.6% operating margin. NVIDIA's ~4% equity holding makes the relationship closer to partnership than rivalry. Its forward P/E is the least stable figure in this table across data vendors (23.1x vs 45.5x on a second sweep).

COMPANY	TICKER	MARKET C...	REVENUE T...	REV GROW...	GROSS MA...	EBIT MAR...	AGENT_NOT...
Qualcomm	QCOM	177.26	44.07	2.0%	54.0%	24.0%	In NVIDIA's 10-K peer group but competing in a different arena — edge and on-device inference, a separate ~\$25-35B market. 1.9% growth is the slowest in the set.
Marvell	MRVL	196.04	9.45	n/a	62.0%	19.0%	OUT OF MARKET-CAP ORDER DELIBERATELY, and excluded from the growth median and the scatter chart below: the only growth figure available is +8.4% QUARTER-ON-QUARTER, which is not comparable to the TTM basis of this column. Data centre is 79% of revenue and grew 46% YoY. Most expensive multiple in the set. NVIDIA took a stake in it in March 2026.

COMPANY	TICKER	MARKET C...	REVENUE T...	REV GROW...	GROSS MA...	EBIT MAR...	AGENT_NOT...
Peer median (6 peers, excl. NVIDIA)				40.0%	58.0%	21.0%	Medians over the six companies named in NVIDIA's own 10-K competitor group (AMD, AVGO, INTC, MRVL, MU, QCOM) — one membership criterion, applied consistently. TSMC fails it (it is NVIDIA's foundry, not a named competitor) and is a memo row only. The growth median uses 5 values because Marvell has no comparable TTM figure. The EV/EBITDA median 17.7x is the median of MU 11.2, QCOM 13.1, INTC 15.6, AVGO 19.8, AMD 34.5, MRVL 40.2.

COMPANY	TICKER	MARKET C...	REVENUE T...	REV GROW...	GROSS MA...	EBIT MAR...	AGENT_NOT...
NVIDIA vs peer median (%)				111.0%	28.0%	2.04	QUALITY: NVIDIA is 2nd of 7 on gross margin, EBIT margin and FCF margin, and 2nd of 6 on revenue growth (Marvell unranked). PRICE: 24.7x forward P/E is a 12.8% PREMIUM to the 21.9x peer median, and 20.3x forward EV/EBITDA is 14.7% ABOVE the 17.7x peer median. Both rows are premiums — modest ones for the second-highest quality profile in the set, but premiums, and they are described as such.
CORRECTION CARRIED FORWARD — the EV/EBITDA row is a PREMIUM, not 'in line'							An earlier draft of this workbook carried 21.9x as the peer EV/EBITDA median. That figure is the peer P/E MEDIAN copied one row up. The correct EV/EBITDA median of the six non-NVIDIA peers is 17.7x, which reverses the reading of that row from 'in line with peers' to '14.7% ABOVE peers'. Any prose that still says NVIDIA trades in line with peers on EV/EBITDA is wrong and is superseded here.

COMPANY	TICKER	MARKET C...	REVENUE T...	REV GROW...	GROSS MA...	EBIT MAR...	AGENT_NOT...
SOURCE SENSITIVITY — peer forward multiples are NOT reliably obtainable							Two independent market-data sweeps on the SAME date (2026-09-08) return materially different forward P/Es for the same companies: AMD 33.0 vs 43.0, Intel 23.1 vs 45.5, Micron 13.85 vs 6.56, Marvell 50.0 vs 41.0, Qualcomm 18.18 vs 16.75. Only Broadcom agrees (20.7 vs 20.68). The peer median P/E is 21.9x on sweep A and 30.8x on sweep B. Sweep A is retained as the primary table because it is one internally consistent collection, but no PRECISE premium or discount to the peer group is defensible from it. WHAT SURVIVES BOTH SWEEPS: NVIDIA's economics are top-two on every measure obtainable, and it carries no large multiple premium for them.

COMPANY	TICKER	MARKET C...	REVENUE T...	REV GROW...	GROSS MA...	EBIT MAR...	AGENT_NOT...
MEMO — TSMC, supplier, EXCLUDED from every median	TSM	2,220	155	37.0%	68.0%	60.0%	NVIDIA's foundry, not a named competitor in its 10-K, so excluded from every median by the same single criterion used to include the six peers. Shown because it is the second-largest company in the arena and its pricing directly sets NVIDIA's gross margin ceiling. Including it would move the peer median P/E to 20.7x.

COMPANY	TICKER	MARKET C...	REVENUE T...	REV GROW...	GROSS MA...	EBIT MAR...	AGENT_NOT...
THE FINDING — this peer group pays nothing for quality							A regression of forward P/E on a quality score (TTM revenue growth x TTM EBIT margin) across the five peers with computable scores gives $R^2 = 0.383$ with a NEGATIVE slope; removing Micron, whose memory-cycle peak dominates the fit, collapses R^2 to 0.007 on $n = 4$ (approximately 0.91). At that level the fit carries no sign information: the honest reading is NO relationship between quality and multiple, not a negative one. The same holds on the second sweep. Consequence for the memo: NVIDIA has little multiple cushion left to lose and correspondingly little multiple risk — the downside case has to express itself through EARNINGS, not through de-rating.

COMPANY	TICKER	MARKET C...	REVENUE T...	REV GROW...	GROSS MA...	EBIT MAR...	AGENT_NOT...
DATA GAPS — stated, not filled							(1) No peer ROIC and no peer net margin were obtainable from this corpus, so neither row appears; NVIDIA's own figures are 68.8% and 63.7%. (2) No forward EV/Revenue was obtainable for the peers. (3) Marvell has no comparable TTM growth figure. (4) AMD's gross margin is a single GAAP quarter, not TTM. (5) Peer market data is dated 2026-09-04 to 2026-09-08 and per-company as-of dates differ by up to four days, so the price side of this table is not struck on one date. (6) NVIDIA's 24.7x / 20.3x are FY2027E forward multiples while the peer column is NTM consensus — a real basis difference, disclosed rather than reconciled away.

Appendix II

EVERY SCE...

TIME ANCHOR
— fiscal years
end in late
January;
FY2026A is the
year ended
January 2026
and FY2027E—
FY2031E are
the model's
five explicit
forecast years.
Model figures
as of 2026-09-
08. Current
price \$230.36
is the 2026-
09-04 close,
re-confirmed
2026-09-08.
UNITS — \$M
unless a row
label says
(\$/share), (%),
(x) or (M).

Line item (\$M unless the label says otherwise)	FY2026A	FY2027E	FY2028E	FY2029E	FY2030E	FY2031E	Source / note
--	---------	---------	---------	---------	---------	---------	---------------

SECTION 1 — KEY ASSUMPTIONS (Base Case)

Consolidated revenue (\$M)	215,909	374,401	474,927	553,064	612,387	654,447	Revenue_Mode l row 7. FY2026A ties to NVDA_Historic als/Income_St atement.
Revenue growth (% y/y)		73.0%	27.0%	16.0%	11.0%	7.0%	DERIVED, never assumed: the mix-weighted result of five segment growth rates.
Data Center revenue growth (% y/y)		80.0%	28.0%	17.0%	11.0%	7.0%	Assumptions row 3. 89.7% of FY2026 revenue. Decelerates to the 15.3% market CAGR.
Gross margin (%)	71.0%	74.0%	72.0%	72.0%	72.0%	72.0%	74.0% in FY2027, then management's guided 72– 73% floor from FY2028.
EBIT margin (%, emergent)	60.0%	63.0%	62.0%	62.0%	62.0%	62.0%	OUTPUT, not an input: gross profit less opex, where opex holds at the FY2026A 10.69% of revenue.

EVERY SCE...

Cash tax rate on EBIT (%)	15.0%	15.0%	16.0%	17.0%	17.0%	17.0%	FY2026 effective rate 15.1%, drifting to 17.0% on OECD Pillar Two.
D&A (% of revenue)	1.0%	1.0%	1.0%	2.0%	2.0%	2.0%	Rising as the installed base ages; NVIDIA is fabless so the absolute figure stays small.
CapEx (% of revenue)	3.0%	3.0%	3.0%	2.0%	2.0%	2.0%	FY2026A actual 2.8%; fabless model keeps capex light.
Working-capital draw (% of INCREMENTAL revenue)	31.0%	22.0%	22.0%	22.0%	22.0%	22.0%	FY2026A actual was 31%; normalised to 22% as growth decelerates. Largest early-year drag on the bridge.
WACC (%)	12.0%						WACC_Model B14. Book weights: equity 94.9% / gross debt 5.1%; cost of equity 12.0% (4.0% risk-free + 1.60 beta x 5.0% ERP).
Terminal growth, perpetual (%)	3.0%						Long-run nominal GDP; 8.6pp below WACC.
Diluted shares (M)	24,100						NVDA_Historicals/Key_Metrics. No buyback effect modelled — conservative against \$55,315M of TTM repurchases.
SECTION 1b — THE CENTRAL ASSUMPTION: FY2028 REVENUE IS THE WHOLE DISAGREEMENT							
Estimate	This workbook (Base)	Consensus (midpoint)	Management guide	Delta vs consensus (%)	Delta vs guidance (pp)		Source / note
FY2028 revenue (\$M)	474,927	681,500	~70% growth (no \$ guide)	-0.30	—		THE ROW THAT CARRIES THE VALUATION. Consensus range \$673,000M–\$690,000M. This single row explains most of the gap between \$139.24 and \$230.36.

EVERY SCE...

FY2028 revenue growth (% y/y)	27.0%	—	70.0%	—	-43.0%	Management issued an unusual long-range ~70% FY2028 growth guide at the Q2 FY2027 print (2026-08-26). The model has 26.8% — a 43.2 percentage-point contradiction of company guidance, not a modelling nuance.
FY2027 revenue (\$M)	374,401	401,270	—	-0.07	—	Consensus range \$391,300M–\$411,240M. The headline –6.7% is DILUTED by H1 actuals of \$177,800M.
H2 FY2027 revenue — the unreported half (\$M)	196,601	228,000	—	-0.14	—	Consensus = the company-guided Q3 \$108,000M plus a ~\$120,000M Q4. On the forecast portion alone the model is 13.8% below, not 6.7%.
Implied Q3 FY2027 revenue (\$M)	97,600	—	108,000	—	-0.10	DERIVED, not modelled — the model is annual. \$196,601M of H2 at a constant ~1.45% sequential rate off Q2's \$96,200M. 9.6% below the company's own \$108.0B +/-2% guide for a quarter already one-third elapsed. The sharpest single test of the FY2027 path, and it is not favourable to us.

EVERY SCE...

HONESTY ROW
 — the model
 was built
 PRICE-BLIND
 and
 GUIDANCE-
 BLIND. The
 projections
 bag never read
 a share price
 and never saw
 the ~70%
 FY2028 guide;
 it was built
 from the sector
 bag's share-
 erosion
 evidence and a
 15.3% market
 CAGR. This
 workbook's
 own view is
 that FY2028
 revenue is THE
 MOST LIKELY
 ERROR IN THIS
 MODEL. Part of
 the 40.6% gap
 is therefore
 this model
 being
 conservative
 rather than the
 market being
 wrong.

SECTION 2 —
 FREE CASH
 FLOW BRIDGE
 (Base Case,
 unlevered)

NOPAT (\$M)	110,699	200,300	246,593	285,454	314,180	335,758	EBIT taxed directly (unlevered). Excludes the non-operating investment gains that inflate reported net income.
Add: D&A (\$M)	2,843	5,242	7,124	8,849	10,411	11,780	Revenue x the D&A assumption (row 11).
Less: capital expenditure (\$M)	-6,042	-9,734	-11,873	-13,274	-14,085	-14,398	Revenue x the capex assumption (row 12). FY2026A actual \$6,042M.
Less: increase in working capital (\$M)	-26,578	-34,868	-22,116	-17,190	-13,051	-9,253	22% of INCREMENTAL revenue from FY2027E; FY2026A actual draw was 31%.

EVERY SCE...

FREE CASH FLOW (\$M)	80,922	160,939	219,728	263,839	297,454	323,887
----------------------	--------	---------	---------	---------	---------	---------

FY2026A restated on the SAME unlevered basis as the forecast. NVIDIA's REPORTED FY2026 FCF was \$96,676M (CFO \$102,718M less capex) — a levered definition; the two must not be compared directly.

FCF margin (%)	37.0%	43.0%	46.0%	48.0%	49.0%	49.0%
----------------	-------	-------	-------	-------	-------	-------

Unlevered basis in every column, so the trend is like-for-like.

Appendix III

Valuation Multiples

MULTIPLE	NVIDIA NOW (X)	NVIDIA 5-YR A...	PEER MEDIAN, ...	GAP VS BENCH...	NOTE
SECTION 1 — WHERE NVIDIA TRADES NOW, AGAINST ITS OWN HISTORY AND AGAINST ITS PEERS. All figures are multiples (x), struck 2026-09-08 at a \$230.36 share price and a \$5,551,676M market capitalisation.					
Forward P/E (FY2027E)	24.7	not sourced	21.9	13.0%	A 12.8% PREMIUM to the 21.9x peer median, not 'in line' with it. It is a modest premium for the second-highest quality score in the peer set, but it is a premium and is named as one. Benchmark: peer median.
Forward P/E (FY2028E, on consensus earnings)	14.8	not sourced	n/a		Falls to 14.8x on consensus FY2028 earnings — but only if consensus FY2028 is right, which is precisely what this workbook's model disputes (our FY2028 revenue is \$474,927M against consensus \$681,500M).

MULTIPLE	NVIDIA NOW (X)	NVIDIA 5-YR A...	PEER MEDIAN, ...	GAP VS BENCH...	NOTE
Forward EV/EBITDA (FY2027E)	20.3	45.9	17.7	15.0%	THE HEADLINE, WITH BOTH BASES STATED. 20.3x against NVIDIA's own five-year average of 45.9x — and BELOW its own five-year low of 40.0x. But the peer comparison moves the OTHER WAY: 20.3x is 14.7% ABOVE the 17.7x peer EV/EBITDA median. Benchmark for the gap column: peer median. See Section 3 — an early draft carried 21.9x in the peer-median cell here, which is the peer P/E median from row 3 copied one row up.
EV/Sales (FY2027E)	13.4	not sourced	n/a		High in absolute terms, and appropriate for a business carrying a 65.2% TTM EBIT margin.
Price / Book (current)	31.8	not sourced	n/a		Book value is a poor anchor for a fabless designer whose assets are largely intangible.
SECTION 1B — EV/EBITDA HISTORY BAND (x). This is the ONLY multiple in the table with a sourced five-year history; see the gap note in Section 3. History spans FY2022-FY2026 (Finbox provider series).					
EV/EBITDA 5-yr high (Jan 2024)	68.2				The peak of the AI re-rating.
EV/EBITDA 5-yr average (FY2022-FY2026)	45.9				The five-year mean of the provider series.
EV/EBITDA 5-yr low (Jan 2025)	40				The bottom of the five-year range — and still nearly double where the stock trades today.

MULTIPLE	NVIDIA NOW (X)	NVIDIA 5-YR A...	PEER MEDIAN, ...	GAP VS BENCH...	NOTE
EV/EBITDA forward FY2027E (today)	20.3				Today's forward multiple sits BELOW the five-year LOW. The stock has de-rated out of its own historical range entirely.
Compression vs own 5-yr average (%)				-56.0%	A 56% compression in EV/EBITDA while revenue grew 83% year over year. BASIS IS MIXED: 20.3x is a FORWARD (FY2027E) multiple and the 45.9x history is a provider series of unstated, probably TRAILING, basis. Read -55.8% as an UPPER BOUND on true compression, not a like-for-like measurement. Benchmark: NVIDIA's own five-year average. The conclusion that today's multiple sits below the five-year low survives on either basis; the magnitude does not.

SECTION 2 —
PEER SET. Forward P/E and EV/EBITDA are NTM (next twelve months); P/S is TTM; quality score is TTM. Peer market data struck 2026-09-04 to 2026-09-08.

Company	Forward P/E, NTM (x)	EV/EBITDA, NTM (x)	P/S, TTM (x)	Quality score (unitless)	Note
Marvell	50	40.2	20.74	n/a	Most expensive multiple in the set. Quality score not computable — no comparable TTM growth figure.
AMD	33	34.5	18.39	6.0%	The clearest mispricing in the table relative to NVIDIA: 8.8x LOWER quality score at a 50% HIGHER forward multiple.

MULTIPLE	NVIDIA NOW (X)	NVIDIA 5-YR A...	PEER MEDIAN, ...	GAP VS BENCH...	NOTE
Intel	23.1	15.6	8.83	0.6%	Near-zero quality score at a 23.1x forward multiple — a turnaround option premium, not an earnings multiple.
NVIDIA	22	n/a	18.26	54.0%	4th of 7 on multiple, 2nd of 7 on quality score. NOTE THE TWO NVDA P/E FIGURES: 22.0x here is an NTM multiple struck on a different date from the 24.7x FY2027E figure in row 3. Both are shown rather than reconciled into one. No forward EV/EBITDA was obtainable for NVIDIA on the peer basis.
Broadcom	20.7	19.8	19.08	24.0%	The closest true comparable: custom-silicon exposure, 48.6% EBIT margin. NVIDIA carries 2.3x Broadcom's quality score at a 6% higher multiple.
Qualcomm	18.18	13.1	4.02	0.4%	Cheapest on sales in the set at 4.0x, reflecting 1.9% growth. Different arena (edge/on-device).
Micron	13.85	11.2	12.74	110.0%	Highest quality score AND cheapest on both P/E and EV/EBITDA — the clearest evidence this peer group does not price quality. Built on a memory up-cycle peak the market plainly expects to mean-revert.

MULTIPLE	NVIDIA NOW (X)	NVIDIA 5-YR A...	PEER MEDIAN, ...	GAP VS BENCH...	NOTE
Peer median, ex-NVDA	21.9	17.7			THE TWO MEDIANS, KEPT APART. 21.9x is the peer P/E median; 17.7x is the peer EV/EBITDA median (Micron 11.2, Qualcomm 13.1, Intel 15.6, Broadcom 19.8, AMD 34.5, Marvell 40.2). Confusing the two is the error corrected in Section 3.

SECTION 3 — THE CORRECTION, AND WHAT IS NOT HERE.

CORRECTION CARRIED FORWARD. An early draft labelled 21.9x as the peer EV/EBITDA median in row 5. It is not — 21.9x is the peer P/E median from row 3, copied one row up. The correct EV/EBITDA median of the six non-NVDA peers is 17.7x. That reverses the reading of the row from 'in line with peers' to '14.7% ABOVE peers'. Both halves of the sentence must travel together: NVIDIA is cheap against its own history and expensive against its peer group on EV/EBITDA.

THE SECOND HALF OF THE SAME CORRECTION. The forward P/E of 24.7x is a 12.8% PREMIUM to the 21.9x peer median, not 'in line' with it, as an earlier version said.

MULTIPLE	NVIDIA NOW (X)	NVIDIA 5-YR A...	PEER MEDIAN, ...	GAP VS BENCH...	NOTE
GAP — NO FIVE-YEAR P/E SERIES. Four of six history cells in Section 1 read 'not sourced'. No corroborated five-year series was found for forward P/E, EV/Sales or P/B, so those cells carry no estimate. 'Current versus history' is therefore delivered for ONE multiple of six, and the history chart below is on EV/EBITDA rather than P/E for that reason.					
GAP — THE PEER MEDIAN INHERITS AN INCONSISTENCY. NVDA_Sector_Analysis states Marvell is excluded from its median, but 21.9x requires INCLUDING it (excluding Marvell gives 20.7x). Carried forward and flagged rather than silently re-derived.					
GAP — THE MAGNITUDE OF THE PEER GAP IS NOT DETERMINABLE. A second independent market-data sweep on the same date (2026-09-08) put the peer P/E median at 30.8x and NVIDIA at a 28.7% DISCOUNT. What is robust across both sweeps is only the DIRECTION on forward earnings. The 12.8% premium figure above is the reading from this table's data and is quoted as such.					

MULTIPLE

NVIDIA NOW (X)

NVIDIA 5-YR A...

PEER MEDIAN, ...

GAP VS BENCH...

NOTE

NO MULTIPLE-BASED PRICE TARGET IS SET HERE, deliberately. This workbook's valuation anchor is the DCF, not a multiple; the sensitivity around that DCF is in Appendix III's Sensitivity_Matrix pod.

Sensitivity Matrix

ANCHORS — ...

Anchor	Value	Where it sits
Base case intrinsic value (USD/share)	139.24	GRID 1, the yellow cell at D16 — row 11.6% WACC (11.581%, read live from WACC_Model B14) x column 3.0% terminal growth. That cell ties to NVDA_Financial_Projections/DCF_Model at \$139.24386168 exactly, to the last decimal.
Bear case intrinsic value (USD/share)	78.83	Not a grid cell — the Bear case moves revenue and margin together, which no single grid axis does. Weighted 25%.
Bull case intrinsic value (USD/share)	189.8	Not a grid cell, for the same reason. Weighted 25%. Note that even the Bull case sits below the market price.
Probability-weighted intrinsic value (USD/share)	136.78	Bear 25% / Base 50% / Bull 25%. 40.6% below the market price.
Current market price (USD/share)	230.36	OFF BOTH GRIDS. No cell in either grid reaches it. Grid 1's most generous corner is \$207.84; Grid 2's is \$226.47.

ANCHORS — ...

Terminal growth the market price implies (%)	7.0%	REVERSE-SOLVED at the model's 11.58% WACC, holding FY2027-31 cash flows fixed: 6.95% PERPETUAL growth, against the 3.0% assumed and the 4.0% right-hand edge of Grid 1. A perpetuity growing at 6.95% requires NVIDIA to outgrow the economy forever.				
Five-year revenue CAGR the market price implies (%)	44.0%	REVERSE-SOLVED at 3.0% terminal growth and a 61.8% terminal EBIT margin: 44.2%, against 24.8% in the Base path, 31.3% in the Bull case and the 37.0% bottom row of Grid 2 — i.e. FY2031 revenue of about \$1.35 trillion.				
Most generous cell in Grid 1 (USD/share)	207.84	9.6% WACC x 4.0% terminal growth — still \$22.52 short of the market price.				
Most generous cell in Grid 2 (USD/share)	226.47	37.0% revenue CAGR x 74.0% terminal EBIT margin — the closest any cell gets, and still \$3.89 short of the market price.				
GRID 1 — IMPLIED INTRINSIC VALUE PER SHARE (USD), WACC (rows) x TERMINAL GROWTH (columns). Time anchor: the same FY2027E-FY2031E cash flows as the Base case, discounted mid-year, terminal value at the year-4.5 factor. This grid holds the business constant and varies only the discount rate and the perpetuity.						
WACC (%) \ terminal growth (%)	2.0%	3.0%	3.0%	4.0%	4.0%	Note

ANCHORS — ...

0.10	162.01	171.05	181.45	193.56	207.84	Most generous WACC row. Even here the grid tops out at \$207.84 — below the \$230.36 price.
------	--------	--------	--------	--------	--------	---

0.11	142.94	149.72	157.39	166.14	176.22	
------	--------	--------	--------	--------	--------	--

0.12	128.11	133.37	139.24	145.85	153.32	THE BASE ROW. WACC 11.581%, read live from NVDA_Financial_Projections/WACC_Model B14. The 3.0% column is the Base case.
------	--------	--------	--------	--------	--------	---

0.13	115.63	119.77	124.35	129.42	135.09	
------	--------	--------	--------	--------	--------	--

0.14	105.53	108.87	112.52	116.54	120.98	
------	--------	--------	--------	--------	--------	--

Every cell in Grid 1 is a FULL DCF recomputation carried out in NVDA_Financial_Projections/Sensitivity_Math, not an interpolation. The grid spans WACC 9.6%-13.6% precisely because three of the WACC inputs — a 4.0% risk-free rate, a 5.0% equity risk premium and a 1.6 beta — are ASSUMPTIONS with no source in this corpus, and a reader should be able to price a different view.

GRID 2 — IMPLIED INTRINSIC VALUE PER SHARE (USD), FIVE-YEAR REVENUE CAGR (rows) x TERMINAL EBIT MARGIN (columns). Time anchor: revenue rebuilt from the FY2026A base of \$215,909M over FY2027E-FY2031E at each row's CAGR, at Base tax, D&A, capex and working capital and a 3.0% terminal growth rate. This grid holds the discount rate constant and varies the business.

revenue CAGR (%) \ terminal EBIT margin (%)	50.0%	56.0%	62.0%	68.0%	74.0%	Note
--	-------	-------	-------	-------	-------	------

0.13	67.1	75.62	83.86	92.67	101.2	
------	------	-------	-------	-------	-------	--

ANCHORS — ...						
0.19	81.98	92.72	103.1	114.19	124.93	
0.25	99.07	112.4	125.29	139.07	152.41	THE BASE-PATH ROW — the CAGR the Base case actually runs. Its centre cell is \$125.29 and is NOT the \$139.24 Base case; see the note below.
0.31	120.37	137	153.07	170.26	186.88	Roughly the Bull case's 31.3% CAGR.
0.37	144.65	165.11	184.88	206.01	226.47	The bottom row is still 7.2 percentage points short of the 44.2% CAGR the market price requires.
GRID 2'S CENTRE IS DELIBERATELY NOT THE BASE CASE, and it is a mistake to read it as one. The centre cell (\$125.29, 24.8% CAGR x 61.8% terminal EBIT margin) holds tax, D&A, capex and working capital at fixed Base averages and assumes a SMOOTH CAGR. The Base case is heavily front-loaded — 73.4% revenue growth in FY2027 — and front-loaded cash flows are worth more under mid-year discounting. That is the whole of the \$13.95 difference between \$139.24 and \$125.29.						

ANCHORS — ...

WHAT IT TAKES
TO REACH
\$230.36, STATED
TWO WAYS. (1)
Hold the
business as
modelled and the
discount rate has
to do the work: at
the model's
11.58% WACC
the price needs
6.95%
PERPETUAL
growth — off the
right-hand edge
of Grid 1, which
stops at 4.0%. (2)
Hold the
discount rate and
the perpetuity
and the business
has to do the
work: the price
needs a 44.2%
five-year revenue
CAGR — off the
bottom of Grid 2,
which stops at
37.0%. Neither
lever reaches the
price inside the
range this model
is willing to draw.

GRID 2 IS THE
WIDER OF THE
TWO, spanning
roughly \$67 to
\$226 per share
against Grid 1's
\$106 to \$208.
That is the
honest result and
it is the point of
running both:
getting the
BUSINESS wrong
costs more than
getting the
DISCOUNT RATE
wrong.

ANCHORS — ...

GAP, DISCLOSED. Neither grid prices the Bear or Bull case, because each of those moves revenue and margin together and no single grid axis does. Neither grid varies the FY2028 revenue assumption directly, which is where the central disagreement with company guidance actually sits — Grid 2's CAGR rows are the nearest proxy. And the WACC used is built on BOOK equity weights, deliberately, to keep the projections model price-independent; market weights would raise WACC by roughly 40bps and lower every cell in both grids.

Reading The Multiples And Grids

The two tables above answer the same question from opposite ends and they do not contradict each other: NVIDIA is cheap against what it used to cost, expensive against what its peers cost, and unreachable at \$230.36 by either lever the discounted cash flow model is willing to pull.

The multiple has already de-rated, and that is the strongest fact in this appendix. Forward EV/EBITDA of 20.3x sits below NVIDIA's own five-year LOW of 40.0x and 56% below the 45.9x five-year average. The stock has left its own historical range entirely — while revenue grew 83% year over year. Two cautions travel with that sentence and must not be dropped. First, the basis is mixed: 20.3x is a forward multiple and the 45.9x / 68.2x / 40.0x history is a single-provider series of unstated, probably trailing, basis, so 56% is an upper bound on the compression rather than a like-for-like measurement. The conclusion that today's multiple is below the five-year low survives either basis; the magnitude does not. Second, no corroborated five-year series exists for forward P/E, EV/Sales or price-to-book, so "current versus history" is delivered for one multiple of six.

Against peers the sign flips, and an early draft got this wrong. That draft labelled 21.9x as the peer EV/EBITDA median. It is not — 21.9x is the peer *P/E* median, copied one row up. The correct EV/EBITDA median across the six non-NVDA peers is 17.7x (Micron 11.2, Qualcomm 13.1, Intel 15.6, Broadcom 19.8, AMD 34.5, Marvell 40.2). NVIDIA's 20.3x is therefore **14.7% ABOVE** peers on that row, not in line with them. The same correction applies on earnings: the 24.7x forward P/E (FY2027E) is a **12.8% premium** to the 21.9x peer median, not "in line" with it — a modest premium for the second-highest quality score in the set, but a premium, and now named as one. The magnitude of any peer gap is not determinable in any case: a second market-data sweep on the same date put the peer P/E median at 30.8x. Only the direction is robust.

The grids say the price is out of reach from both directions. The base case of \$139.24 sits at the centre of Grid 1, where an 11.58% WACC meets 3.0% terminal growth, and it ties to the DCF model to the last decimal. Hold the business as modelled and let the discount rate do the work, and the most generous corner the grid will draw — 9.6% WACC at 4.0% terminal growth — reaches \$207.84, still \$22.52 short. Reverse-solved, \$230.36 needs **6.95% perpetual growth**: a company outgrowing the economy forever, and off the right-hand edge of the grid. Hold the discount rate instead and make the business do the work, and Grid 2's bottom-right cell — a 37.0% five-year revenue CAGR at a 74.0% terminal EBIT margin — reaches \$226.47, \$3.89 short. Reverse-solved that way, the price needs a **44.2% five-year revenue CAGR**, against 24.8% in the base path, and implies FY2031 revenue near \$1.35 trillion.

Two things not to misread in Grid 2. Its centre cell is \$125.29, not \$139.24, and that is deliberate rather than an inconsistency: the grid assumes a smooth CAGR while the base case is front-loaded at 73.4% FY2027 revenue growth, and front-loaded cash flows are worth more under mid-year discounting. And Grid 2 is the wider grid — roughly \$67 to \$226 against Grid 1's \$106 to \$208 — which is the honest result of running both. Getting the business wrong costs more than getting the discount rate wrong.

What neither table settles. Neither grid varies the FY2028 revenue assumption directly, and that is where the workbook's real disagreement lives — \$474,927M modelled against \$681,500M consensus and management's guided ~70% growth. Grid 2's CAGR rows are the nearest proxy, not a substitute. Three WACC inputs — a 4.0% risk-free rate, a 5.0% equity risk premium and a 1.6 beta — are assumptions with no source in this corpus, which is precisely why Grid 1 runs from 9.6% to 13.6% and invites a reader to substitute their own.

Frequently Asked Questions

Topic: Valuation

Because the disagreement is with management, not the market. Every modelled scenario sits below \$230.36 - even the bull at \$189.80 - but the -40.6% gap rests on our FY2028 revenue of \$474,927M, 26.8% growth against the roughly 70% management guided. That is a bet against the company's own number, not against exuberance: forward EV/EBITDA of 20.3x is below NVDA's five-year low of 40.0x, though 14.7% above the 17.7x peer median. It resolves on 17 November 2026. So: Hold, low conviction, NO position initiated - the Neutral-spring starter rule is deliberately overridden because the bag's probability-weighted 12-month return is -13.1% against a +2.7% ceiling. An existing holder runs a \$195.00 stop on two consecutive weekly closes, -15.3%.

Sources: Implied vs Current · Our Estimates vs Consensus · Scenario Analysis · Execution Plan

Q2. EVERYONE SAYS NVIDIA IS EXPENSIVE. IS IT?

Topic: Valuation

Not against its own peer group. Forward P/E is 22.0x against a six-peer median of 21.9x, a premium of 0.5%, while NVIDIA ranks 2nd of 7 on gross margin, EBIT margin, FCF margin and revenue growth simultaneously. Against AMD it carries 8.8 times the quality score at a 33% lower multiple. The important caveat is that this peer group does not price quality at all: regressing forward P/E on quality score gives an R-squared of 0.007 excluding Micron, with a negative slope. So the multiple is cheap relative to quality, but quality is not what sets multiples here.

Sources: Peer Valuation Comparison · Peer Group Pays Nothing For Quality NVIDIA Trades At Median

Q3. MANAGEMENT GUIDED 70% GROWTH FOR FY2028. WHAT IS THE SINGLE EARLIEST SIGNAL THAT IT WILL NOT HOLD?

Topic: Thesis Risk

The disappearance of one phrase from the earnings call. The FY2028 guide is explicitly supply-constrained, with management saying customer forecasts point to growth doubling and the 70% figure reflecting what NVIDIA can build. So long as supply binds, demand softness is invisible in reported revenue. The first quarter in which management stops describing the outlook as supply-constrained is the moment the demand-supply gap has closed, and it precedes any deterioration in the numbers. Watch that sentence before watching bookings, backlog or the \$2 trillion cloud backlog figure.

Sources: Supply constraint framing · Risk 6 watch-signal · Q2 FY2027

Q4. WHY DOES YOUR MODEL CARRY NO FURTHER EBIT MARGIN EXPANSION FOR A COMPANY ALREADY EARNING A 60.4% MARGIN IN A BOOMING MARKET?

Topic: Thesis Risk

Note first that the margin does not fall. The emergent EBIT margin, Income_Forecast row 8, runs 60.39% in FY2026A, 63.31% in FY2027E, then 61.81% in FY2028E and flat through FY2031E: it rises, steps down once, and never drops below the FY2026A actual. The 62.0%-to-58.0% path is Assumptions row 10, carried on row 9 for comparison only and used in NO calculation. EBIT here is an output of an explicit opex build at an Operating

Leverage Coefficient of 1.00, so opex holds at 10.69% of revenue. Two thirds of the 27.9-point historical margin gain came from opex falling 27.6% to 9.5% of revenue, a lever that cannot be pulled twice.

Sources: Income Forecast · Key Drivers

Q5. IF CUSTOM ASICS ARE GROWING THREE TIMES FASTER THAN MERCHANT GPUS, HOW LONG BEFORE THEY TAKE NVIDIA'S BUSINESS?

Topic: Thesis Risk

Not soon, and the proximity verdict in this bag is medium-term rather than imminent. Custom ASIC shipments grew 44.6% year over year against 16.1% for merchant GPUs, and ASIC servers should be 27.8% of accelerator server shipments in 2026 — real displacement, but confined to inference and bounded by CUDA on training. NVIDIA's counter is to stop fighting for the die: it took a stake in Marvell in March 2026 and put \$3.5B into MediaTek, opening NVLink Fusion so third-party accelerators still sit in an NVIDIA rack. The falsifier is a frontier training run completing on non-NVIDIA silicon.

Sources: Disruptor Analysis · Competitive Dynamics

Q6. EXPORT CONTROLS CLOSED CHINA. HOW MUCH DAMAGE IS STILL TO COME?

Topic: Thesis Risk

Bounded, and smaller than the coverage implies. China fell to 9.1% of FY2026 revenue from 13.1%, even though China revenue GREW 15.0% in absolute terms — the denominator simply grew faster. The April 2025 H20 licence requirement already produced a \$4.5B charge, and licences granted from February 2026 for H200 have produced no revenue to date. The open question is direction: BIS moved to case-by-case review on 15 January 2026, but reporting points to a new rule targeting Chinese access to remote AI servers, which would widen controls from what NVIDIA ships to how its installed base is used. This workbook records that direction as genuinely undetermined.

Sources: Policy Regulatory · Macro Dashboard

Q7. WHY IS THERE NO TAIWAN SCENARIO WHEN THE WORKBOOK CALLS IT THE ONLY EXISTENTIAL RISK?

Topic: Thesis Risk

Because a five-year DCF is the wrong instrument for it. Lexo's Bear Route 2 is a discrete supply interruption — the business is unchanged, it simply cannot ship — whereas a DCF scenario is a growth-and-margin path. Forcing it into the table would imply a precision about timing and duration that nobody has, and would blend a probability-weighted event into an operating forecast. The modelled Bear case is Lexo's Route 1, the competitive route, where a portable compilation layer reaches training parity and the capex cycle turns; that IS a margin path and can be modelled honestly. The Taiwan exposure is carried where it belongs, as the sole existential_threat on the market lightcone circle.

Sources: Assumptions Narrative · Geopolitical

Q8. AMD KEEPS CLOSING THE GAP ON SPECS. WHY IS THIS OVERVIEW NOT MORE WORRIED ABOUT IT?

Topic: Competition

Because the share NVIDIA has lost has not gone to AMD. NVIDIA fell from roughly 92% of AI accelerator revenue in 2023 to 80-85% in 2026, while AMD's Instinct revenue is estimated at \$7-8B against NVIDIA's \$193.7B of FY2026 Data Center revenue. The loss went to customers' own silicon and to Broadcom, whose AI ASIC revenue passed \$20B in FY2025. AMD's MI350X matches B200 on FP8 compute, but NVIDIA delivers 50-55% model FLOPs utilisation versus roughly 45%, and NVLink runs 1.8 TB/s against about 128 GB/s.

Sources: Named competitors · Counter-evidence

Q9. NVIDIA SPENT ABOUT \$20B ON GROQ. WHAT WAS IT ACTUALLY BUYING?

Topic: Competition

Time on its weakest flank. Custom ASIC shipments are growing 44.6% against 16.1% for merchant GPUs, with parity projected by 2027, and the migration concentrates in stable-model inference — the workload a purpose-built part wins. Rather than build a better GPU for it, NVIDIA licensed a different architecture and put a Language Processing Unit on its own annual cadence, with the Groq 3 LPU slated for H2 2026. Whether it ships on schedule and gets a named successor is the test.

Sources: Finding 2 · LPU cadence

Q10. IF THE STOCK IS 2.6% OFF ITS 52-WEEK HIGH, HOW CAN THE SPRING STATE BE NEUTRAL RATHER THAN EXTENDED?

Topic: Catalysts

Because Extended requires all three of its tests and NVDA passes one. It is within 10% of the 52-week high, at 2.6% below \$236.54 — that test passes. But the multiple test fails badly: 20.3x forward EV/EBITDA against a 45.9x five-year average is far BELOW its history, not 15% above it. And the momentum test fails: +15.1% from the June close of \$200.09 is short of the 20% three-month threshold. Coiled fails immediately, since the stock sits 40.4% above its 52-week low. What makes this genuinely interesting rather than a shrug is the combination — price near highs on a multiple below its five-year low means the entire recovery came from earnings rather than re-rating.

Sources: Price History Narrative · Multiples Analysis

Q11. THE LEXO WORK CALLS NVIDIA AN N-OF-1 BUSINESS AT 83/100. DOES A HOLD CONTRADICT THAT?

Topic: Classification

No, and the two are answering different questions over different horizons. NVDA_Lexo_Analysis classifies NVIDIA as Diamond in Air, Ocean, Luminous on a 3-5 year intrinsic view, and nothing in the trading bag revises that. This Hold is a 6-12 month tactical call about price and about one testable revenue disagreement. A company can be exceptional and fully priced at the same time; separating those two judgments is exactly why the platform keeps every research and modelling bag price-blind and lets only this one see the market price. If the 17 November print confirms FY2028 near 70%, the tactical view becomes a Buy while the intrinsic verdict is unchanged.

Sources: Market Setup · Lexo Symbology

Q12. WHAT IS THE SINGLE MOST AGGRESSIVE ASSUMPTION IN THIS MODEL?

Topic: Data Quality

Working capital, not growth. In FY2026 the change in working capital consumed \$26,578M, which was 31% of that year's incremental revenue. The model normalises that to 22% of incremental revenue on the argument that the drag eases as growth decelerates — but that still costs \$34.9B in FY2027, more than three times the year's capital expenditure, and if the receivables and inventory build proves structural rather than growth-driven then 22% is too generous. NVDA_Internal_Analysis already flags DSO at 76 days and an accruals ratio doubled to 18.3%. Every dollar of that assumption falls straight through to free cash flow.

Sources: FCF Bridge · Earnings Quality

Mind Map Summary

The whole memo reduces to one shape: an exceptional business, a **position** question rather than a quality question, and a **price** question that resolves on a known date — **17 November 2026**.

NVIDIA – Diamond in Air | Ocean | Luminous

N-of-1, 83/100 · Hold · Low

+-----+-----+-----+		
QUALITY: CLOSED	POSITION: OPEN	PRICE: DATED
(22% of the case)	(54% outside control)	(resolves 2026-11-17)
ROIC 68.8%	fc_01 Share ceded (20)	\$230.36 market
GM 74.7% vs 58.4%	fc_05 Credit-funded demand (13)	\$136.78 PW intrinsic
EBIT 65.2% vs 21.5%	fc_07 Rack-scale lock-in (17)	= -40.6% gap
3 of 4 Lexo axes	fc_08 Self-funded loop (7)	Bull \$189.80 still -17.6%
at their ceiling	Ocean holds ONLY because	
=> Every live argument	of rack-scale, while	FY2028 is the
is about something	unit share falls	whole gap:
ELSE	Mix-not-defeat held at	\$474,927M / 26.8%
	PLAUSIBLE, not Strong	vs guided ~70%
	(networking is not HBM-gated,	The counterparty is
	compute is – the evidence	MANAGEMENT, not
	does not discriminate)	market exuberance

1. Quality is closed, so stop diligencing it. A 68.8% ROIC, a 74.7% TTM gross margin against a 58.4% peer median and a 65.2% EBIT margin against 21.5% put three of four Lexo axes at their ceiling. Only 22% of the factor weight sits inside the business at all, and what remains there is measured rather than open: succession (fc_09, weight 8), the investment book (fc_10, weight 9) and a margin lever already spent (fc_11, weight 5). Diligence spent on unit economics is spent in the wrong place.

2. The position is the live question, and most of it is outside NVIDIA's control. 54% of the weight sits in the external and ecosystem circles. Roughly 11-12 points of accelerator share have gone to hyperscaler ASICs since 2023 while gross margin *rose* from 72.7% to 74.7% — the reading that this is mix rather than defeat is the single most consequential claim in the workbook, and it is deliberately held at **Plausible, not Strong**, because networking is not HBM-gated while compute is, so the evidence cannot discriminate between the two stories. The market axis reads Ocean only because of rack-scale platform lock-in (fc_07, weight 17); remove that and the falling unit share would read differently.

3. The price disagreement is narrow, dated, and against an unusual counterparty. The -40.6% gap does not come from disliking the multiple — forward EV/EBITDA of 20.3x sits *below* NVIDIA's own five-year low of 40.0x. It comes from one row: FY2028 revenue of \$474,927M (26.8% growth) against management's guided ~70% and consensus \$681,500M. The model was built price-blind and guidance-blind; it never saw the guide. This workbook's own view is that **FY2028 is the most likely error in it**. That is why the stance is Hold and not Sell, and why the date matters more than the level.

WHAT THE MAP DELIBERATELY LEAVES OUT

The **Taiwan Strait** is held outside the factor composition entirely. It is treated as a position-sizing question, not a stock-selection one: a discrete supply interruption is not a growth-and-margin path, and a five-year DCF is the wrong instrument for it. Its absence from the weighting is a choice, not an oversight.

WHERE THE MAP IS THIN

Three edges are honestly weak. The self-funded demand loop (fc_08) kept its mechanism but **lost its magnitude claim** when the probe failed, which is why it carries only 7. The three-scenario set has no Steelman or Market-Implied branch upstream, so the left tail is understated — the 10% geopolitical bear route is not modelled as a cash-flow path. And the workbook carries two unreconciled current prices in NVDA_Overview; this memo uses \$230.36 throughout.

Glossary

Fourteen terms this memo assumes. Generic finance vocabulary is deliberately omitted; each entry below is tied to the thesis pillar or risk it actually bears on.

Merchant GPU — a general-purpose accelerator NVIDIA designs and sells to anyone, as opposed to silicon a customer designs for itself. The share-cede argument (fc_01) is measured in merchant-GPU terms: ASIC server shipments grew 44.6% year over year against 16.1% for merchant GPUs.

Custom ASIC (also XPU) — an accelerator a hyperscaler designs for its own workloads rather than buying from NVIDIA. Roughly 11-12 points of accelerator share have moved this way since 2023. This is the single largest external factor in the workbook (fc_01, weight 20).

Hyperscaler — the very large cloud operators that buy NVIDIA at scale. Roughly half of Data Center revenue is four customers, which is why the exit trigger is written as any two of them guiding capex flat or down in the late-January 2027 window.

Rack-scale (GB200 / GB300 NVL72) — a whole rack of GPUs, CPUs and switching sold and liquid-cooled as one integrated unit rather than as boards. Rack-scale lock-in (fc_07, weight 17) is the reason the market axis reads Ocean while unit share falls; it is the load-bearing claim of the ecosystem circle.

CUDA — NVIDIA's proprietary parallel-computing software layer and the libraries built on it. The accumulated code and developer habit around it is the switching cost that makes a specification-level competitor insufficient — the reason AMD closing a spec gap does not, by itself, move share.

NVLink — NVIDIA's high-bandwidth GPU-to-GPU interconnect, the fabric that makes a rack behave as a single accelerator. It is what makes rack-scale a product rather than a bundle.

InfiniBand / Spectrum-X — NVIDIA's two data-centre networking stacks (the former from the Mellanox acquisition, the latter Ethernet-based). Networking compounds at 199% against 92% for the segment — roughly \$39B a year incremental against ~\$26B foregone — and is the core of the mix-not-defeat reading.

HBM (High Bandwidth Memory) — the stacked DRAM sitting beside the compute die, and the industry's binding supply constraint. This term is load-bearing in this memo for one reason: **networking is not HBM-gated while compute is**, which is exactly why the workbook refuses to upgrade the mix-not-defeat reading beyond Plausible — the two explanations predict the same data.

CoWoS (Chip-on-Wafer-on-Substrate) — TSMC's advanced packaging process that binds logic and HBM together. Its capacity, not demand, is what management means when it calls the FY2028 guide supply-constrained.

Blackwell / Rubin — the current and next GPU architecture generations. The cadence matters to the thesis because an annual architecture step is what converts the installed base into recurring replacement demand rather than a one-off build-out.

TPU / Trainium / Instinct (MI300) — customer-designed and rival accelerators: Google's TPU, Amazon's Trainium, AMD's Instinct line. AMD's Instinct revenue is estimated at \$7-8B, which is why the ceded share is attributed to customer silicon rather than to AMD.

Groq / LPU (Language Processing Unit) — the inference-specialist silicon company NVIDIA committed roughly \$20B to, and its chip class. The purchase is read in this memo as buying time on the inference flank, which is where custom ASIC migration concentrates.

H20 and the export-control regime — the US restrictions governing what NVIDIA may sell into China, and the down-specified part built for them. China fell to 9.1% of FY2026 revenue from 13.1% even as China revenue grew 15.0% in absolute terms, which bounds the remaining damage.

Supply-constrained guidance — management's framing that the FY2028 outlook reflects what NVIDIA can build, not what customers will buy. The whole memo turns on this phrase: if it holds on 17 November 2026, the workbook's 26.8% FY2028 growth assumption is the error and 20.3x forward EV/EBITDA is genuinely cheap.

Check 9 — the residual, and why it was not fixed here

Internal QA for the NVDA memo build, run 2026-09-12 against the 10-check hard gate. **Result: 8 PASS • 1 FAIL • 1 WARN.**

#	CHECK	RESULT	EVIDENCE
1	Rating consistency	PASS	Hold on the One_Pager Trading Stance line, the Execution_Plan rating row and its so-what opener, and the Outline Exec_Summary. No surface states Buy/Sell/Short as the call; those words appear only inside named trigger conditions.
2	Bull/Base/Bear reconciliation	PASS	Investment_Case and Scenario_Analysis carry an identical 25 / 50 / 25 (Bear decomposed 15% competitive + 10% geopolitical), summing to exactly 100%. Opp gap, so no reconciliation sentence is required.
3	Stop-loss verbatim	PASS	\$195.00 on two consecutive weekly closes in both the One_Pager "What we'd do" paragraph and Execution_Plan. Same confirmation rule, same -15.3% risk-per-unit.
4	Spring state	PASS	Neutral in Market_Setup and Neutral in Execution_Plan, verbatim. Coiled/Extended appear in Market_Setup only as the rule definition being contrasted, not as the state.
5	Current price	PASS	\$230.36 in Key_Facts, Market_Setup, Scenario_Analysis, Investment_Case, Execution_Plan, DCF_Summary and the One_Pager — zero drift. The other in-range figures are the 52-week high (\$236.54) and the \$199–\$201 shelf, both labelled.
5b	Returns anchored to the as-of price	PASS	Every return, target and stop in the One Pager is stated against \$230.36 with the 2026-09-08 as-of date. Market-cap arithmetic checks: $\$230.36 \times 24,100\text{M diluted} = \$5,551,676\text{M}$.
6	PWR ↔ Rating sign	PASS	12-month PWR is -13.1% , inside the ±15% band, so Hold requires no override. Execution_Plan additionally documents why the Neutral-spring starter rule is overridden rather than followed.

#	CHECK	RESULT	EVIDENCE
7	No Lexo vocabulary in the Cover H1	PASS	H1 reads <i>"Exceptional Business, Fully Priced — The Bear Case Bets Against Management's Own Guidance"</i> — 13 words, a complete sentence, zero forbidden terms.
8	Outline ↔ body reconciliation	PASS	The Outline Exec_Summary two-horizon block matches the One_Pager on Rating (Hold), conviction (Low), position (none), price (\$230.36), probability-weighted intrinsic (\$136.78), gap (−40.6%), verdict (Diamond in Air Ocean Luminous) and tier (N-of-1, 83/100).
9	Loop completeness	FAIL (residual)	Both structural assertions pass: all 11 outline points and all 6 classification axes carry evidence. The third fails — 32 of 49 load-bearing decisions are claimed by no outline point or axis . See the breakdown below.
10	Scenario-set coherence	WARN	audit_scenario_coherence returned audited: false — bag NVDA_Lexo_Analysis has no Scenario_Lens pod, so there is nothing to audit. Per the gate this is a WARN, not a FAIL: coherent: true, problems: [].

Of the 32 unclaimed load-bearing decisions:

- **14 carry** section: "memo" and were logged by this memo build itself. The outline was authored in the prior pipeline step and could not have referenced decisions that did not yet exist. These are a pipeline-order artefact, not a broken chain.
- **18 are research-phase conclusions** (business-overview 2, historical-analysis 3, internal-analysis 3, sector-analysis 3, market-analysis 2, lexo-analysis 2, financial-projections 2, emerging-themes 1). Roughly ten of these map cleanly onto an existing outline point — for example pj1PLiGCXHLg (the networking-attach test, compute falling from 91.5% to 82.4% of Data Center) belongs to sp_03, and e63PcuM6IiwA (71.6% of base-case enterprise value in the terminal value) belongs to sp_09.
- **Two must NOT be linked as supporting evidence:** HcHrttsfilrs (the order-of-magnitude demand-support claim) was retracted when its magnitude probe failed, and 8n0btQcb9fB1 carries a superseded valuation (\$124.22 base on a \$67.74–\$173.73 range, since replaced by \$139.24 on \$78.83–\$189.80).

Why no backfill was attempted. The only read path for the outline (read_workbook_state(sections=["outline"])) returns evidence.sourceCount as an integer, not the underlying {resourceId, quote} citation objects. outline_upsert_point is a REPLACE that drops any ref omitted from the call. Backfilling a decision id would therefore have destroyed the source excerpts on every point touched — trading irrecoverable citations for a link count. The gate forbids inventing references; destroying them is worse. The residual is recorded here instead, per the gate's residual path.

Fix, when the outline is next regenerated: have the outline beat steps pass the full decisions array for each point rather than a single representative id. The current outline links exactly one decision per point (two on sp_04), which is what leaves 32 conclusions unclaimed.

ONE UNRECONCILED FIGURE CARRIED FORWARD, NOT EDITED

Featured FAQ Q2 quotes a forward P/E of **22.0x** against the 21.9x six-peer median (a 0.5% premium). The One_Pager and Appendix III carry **24.7x** against the same median (a 12.8% premium). The FAQ pod is a rendering of the workbook_faqs rows and is reproduced verbatim so the exported PDF and the workspace FAQ tab cannot disagree; rewriting the prose here would have broken that guarantee. Flagged for resolution upstream in NVDA_Sector_Analysis/Peer_Valuation_Comparison.

Gate run 2026-09-12 on workbook F4l0orDVfkCX. All figures as of 2026-09-08. No fix-up rounds were dispatched; the single failure is a recorded residual with the reason above.

Lexo Terminology

A reference guide to the Lexo Investment Framework's vocabulary — for readers new to our research.

Symbology — The Four Axes

Every Lexo classification states a company on four orthogonal axes. Example: “JPMorgan is a Diamond in Water | Ocean | Luminous.”

FOUNDATION intrinsic quality		MEDIUM operating environment	
Clay	developing; moat unproven	Mud	opaque, friction-heavy, ambiguous direction
Gold	proven moat; durable franchise	Water	normal market conditions, market-paced
Diamond	exceptional; irreplaceable, category-defining	Air	tailwinds, abundant supply of opportunity
		Space	structural transformation, frontier territory
REACH market scope		LUMINOSITY value extraction	
River	narrow segment	Brilliant	first light; potential visible, not yet extracted
Lake	defined market	Radiant	growing glow; value visible, partially extracted
Ocean	global / multi-vertical	Luminous	full luminosity; best-in-class monetization

Force Types — How Drivers Are Classified

Every factor on the Light Cone is labelled with a force type describing the nature and direction of the driver.

Headwind	temporary negative force (1–4 quarters)
Tailwind	temporary positive force (1–4 quarters)
Wave +	secular positive structural shift (multi-year)
Wave –	secular negative structural shift (multi-year)
Existential Threat	could end the core business (e.g., Kodak / digital photography; Blockbuster / streaming)
Mud (force)	direction genuinely unclear — used honestly when classification is ambiguous

Note on “Mud”. The word appears in two places with different meanings. As a *Medium state*, “Diamond in Mud” describes the company's *environment* — opaque, friction-heavy. As a *Force Type*, “Mud” on the Light Cone describes a specific driver whose direction is unclear. Context disambiguates.

Color & Trajectory Semantics

- Purple #9B59F5**Intrinsic quality — the Lexo brand color. Foundation tier markers, subject company in schematics, “Our view” callouts.
- Emerald #10b981**Positive momentum — tailwinds, upward trajectories, internal-circle factors on the Light Cone.

Trajectories are stated as direction-of-travel, not just a current snapshot:

Today:	Diamond in Water Ocean Luminous		
3-Year Bull:	Diamond in Air	(if rate cycle turns favorable)	
3-Year Bear:	Diamond in Mud	(if regulatory tightening lands)	
5-Year Potential:	Diamond in Space	(if AI / structural transition succeeds)	

DISCLAIMER: This financial memo is generated by Lexo AI and is intended for informational purposes only. It does not constitute investment advice, a solicitation, or an offer to buy or sell any security. Past performance is not indicative of future results. All investments involve risk, including the possible loss of principal. The information contained herein has been obtained from sources believed to be reliable, but its accuracy and completeness are not guaranteed. Recipients should conduct their own due diligence and consult with a qualified financial advisor before making any investment decisions. Lexo AI and its affiliates disclaim any liability for any direct or indirect loss arising from the use of this document. © September 21, 2026 Lexo Financial Intelligence.